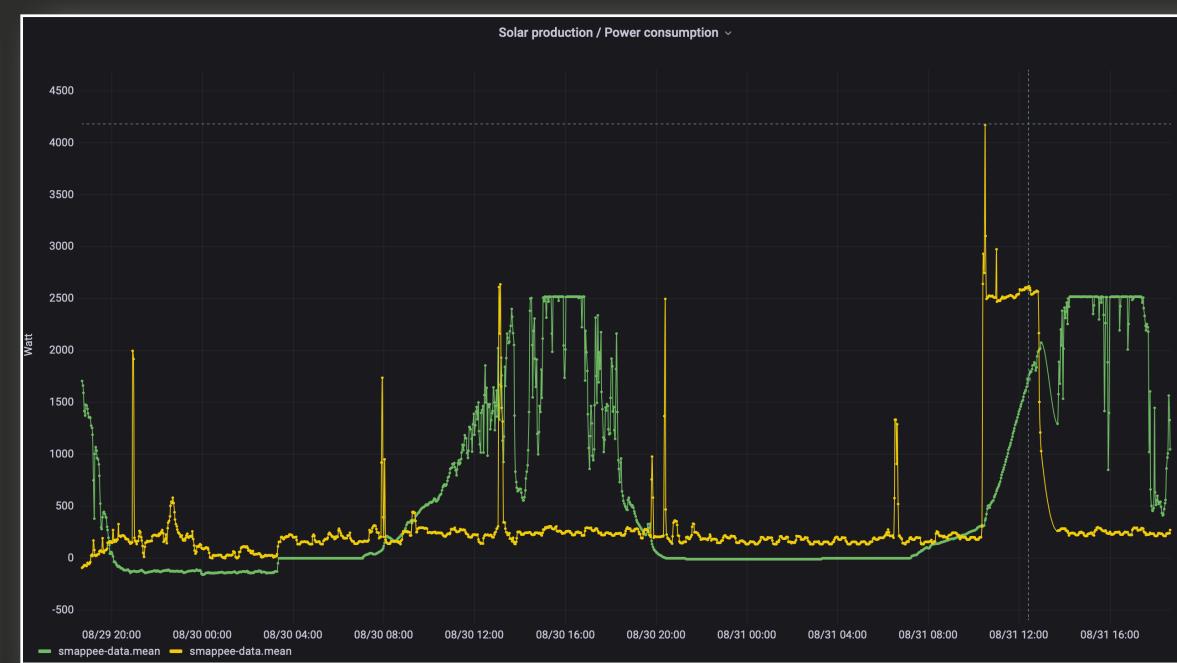
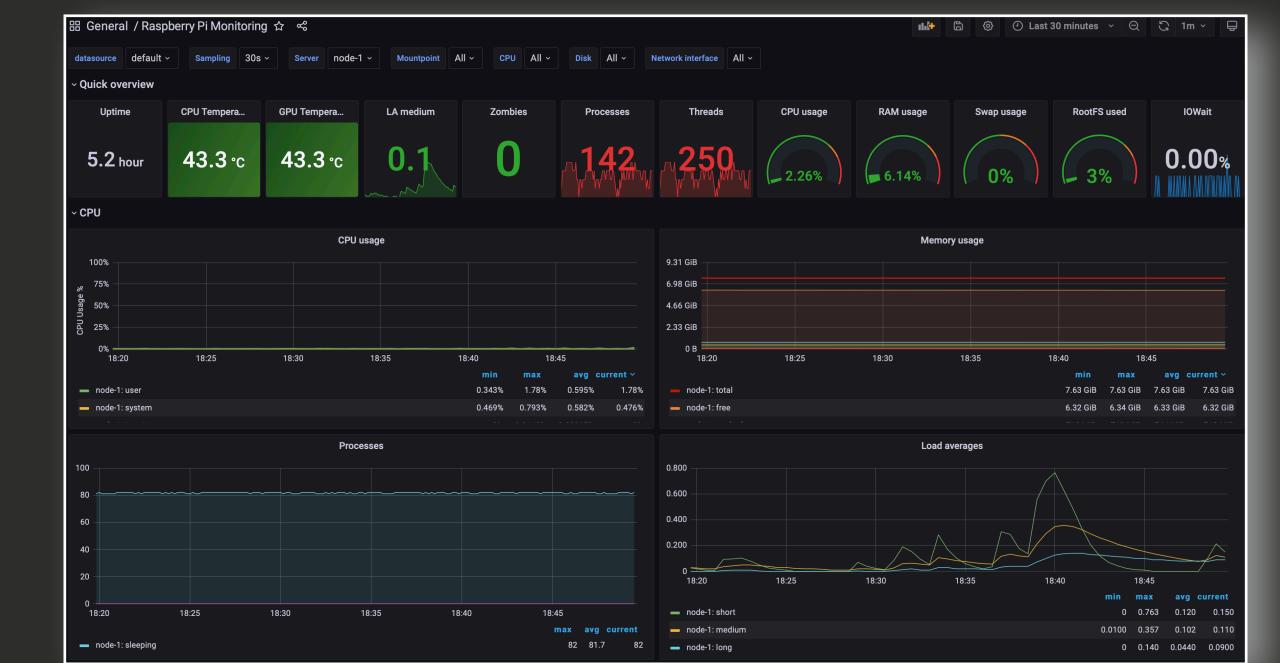


30/09/2022 – A LUNATECH PRESENTATION



Setting up a home energy monitoring system on Raspberry Pi with OSS



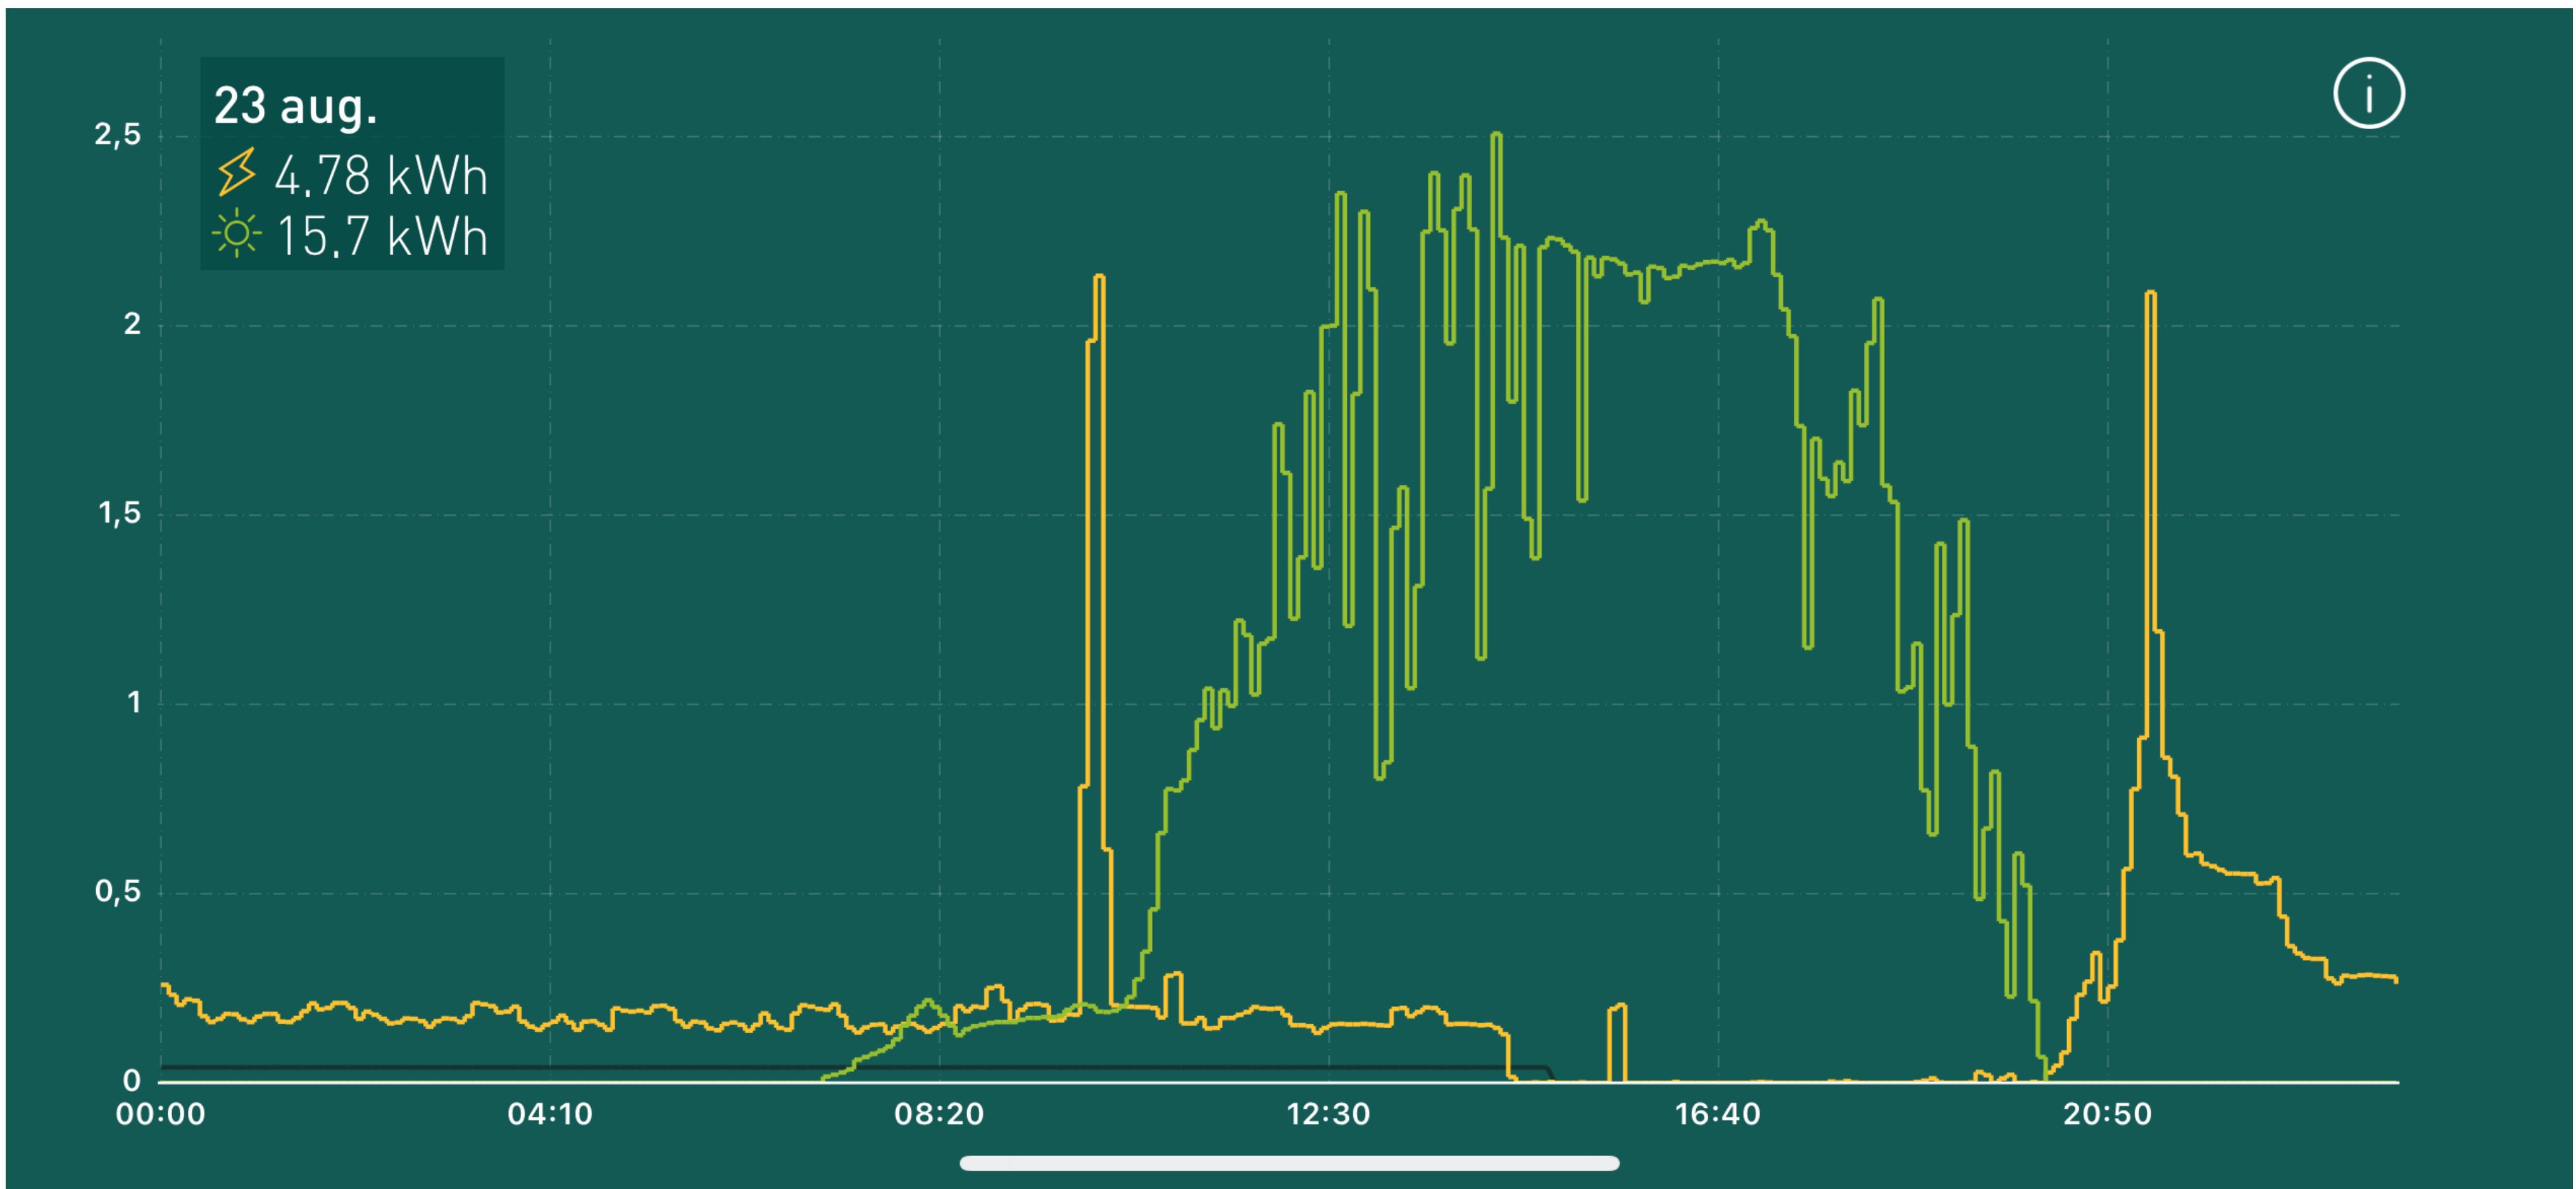
Eric Loots - Lunatech

Agenda

- Why did I start this project?
- Goals of the project
- Choosing the core OSS components
- Practical experience while implementing the project
- Conclusions
- What's next

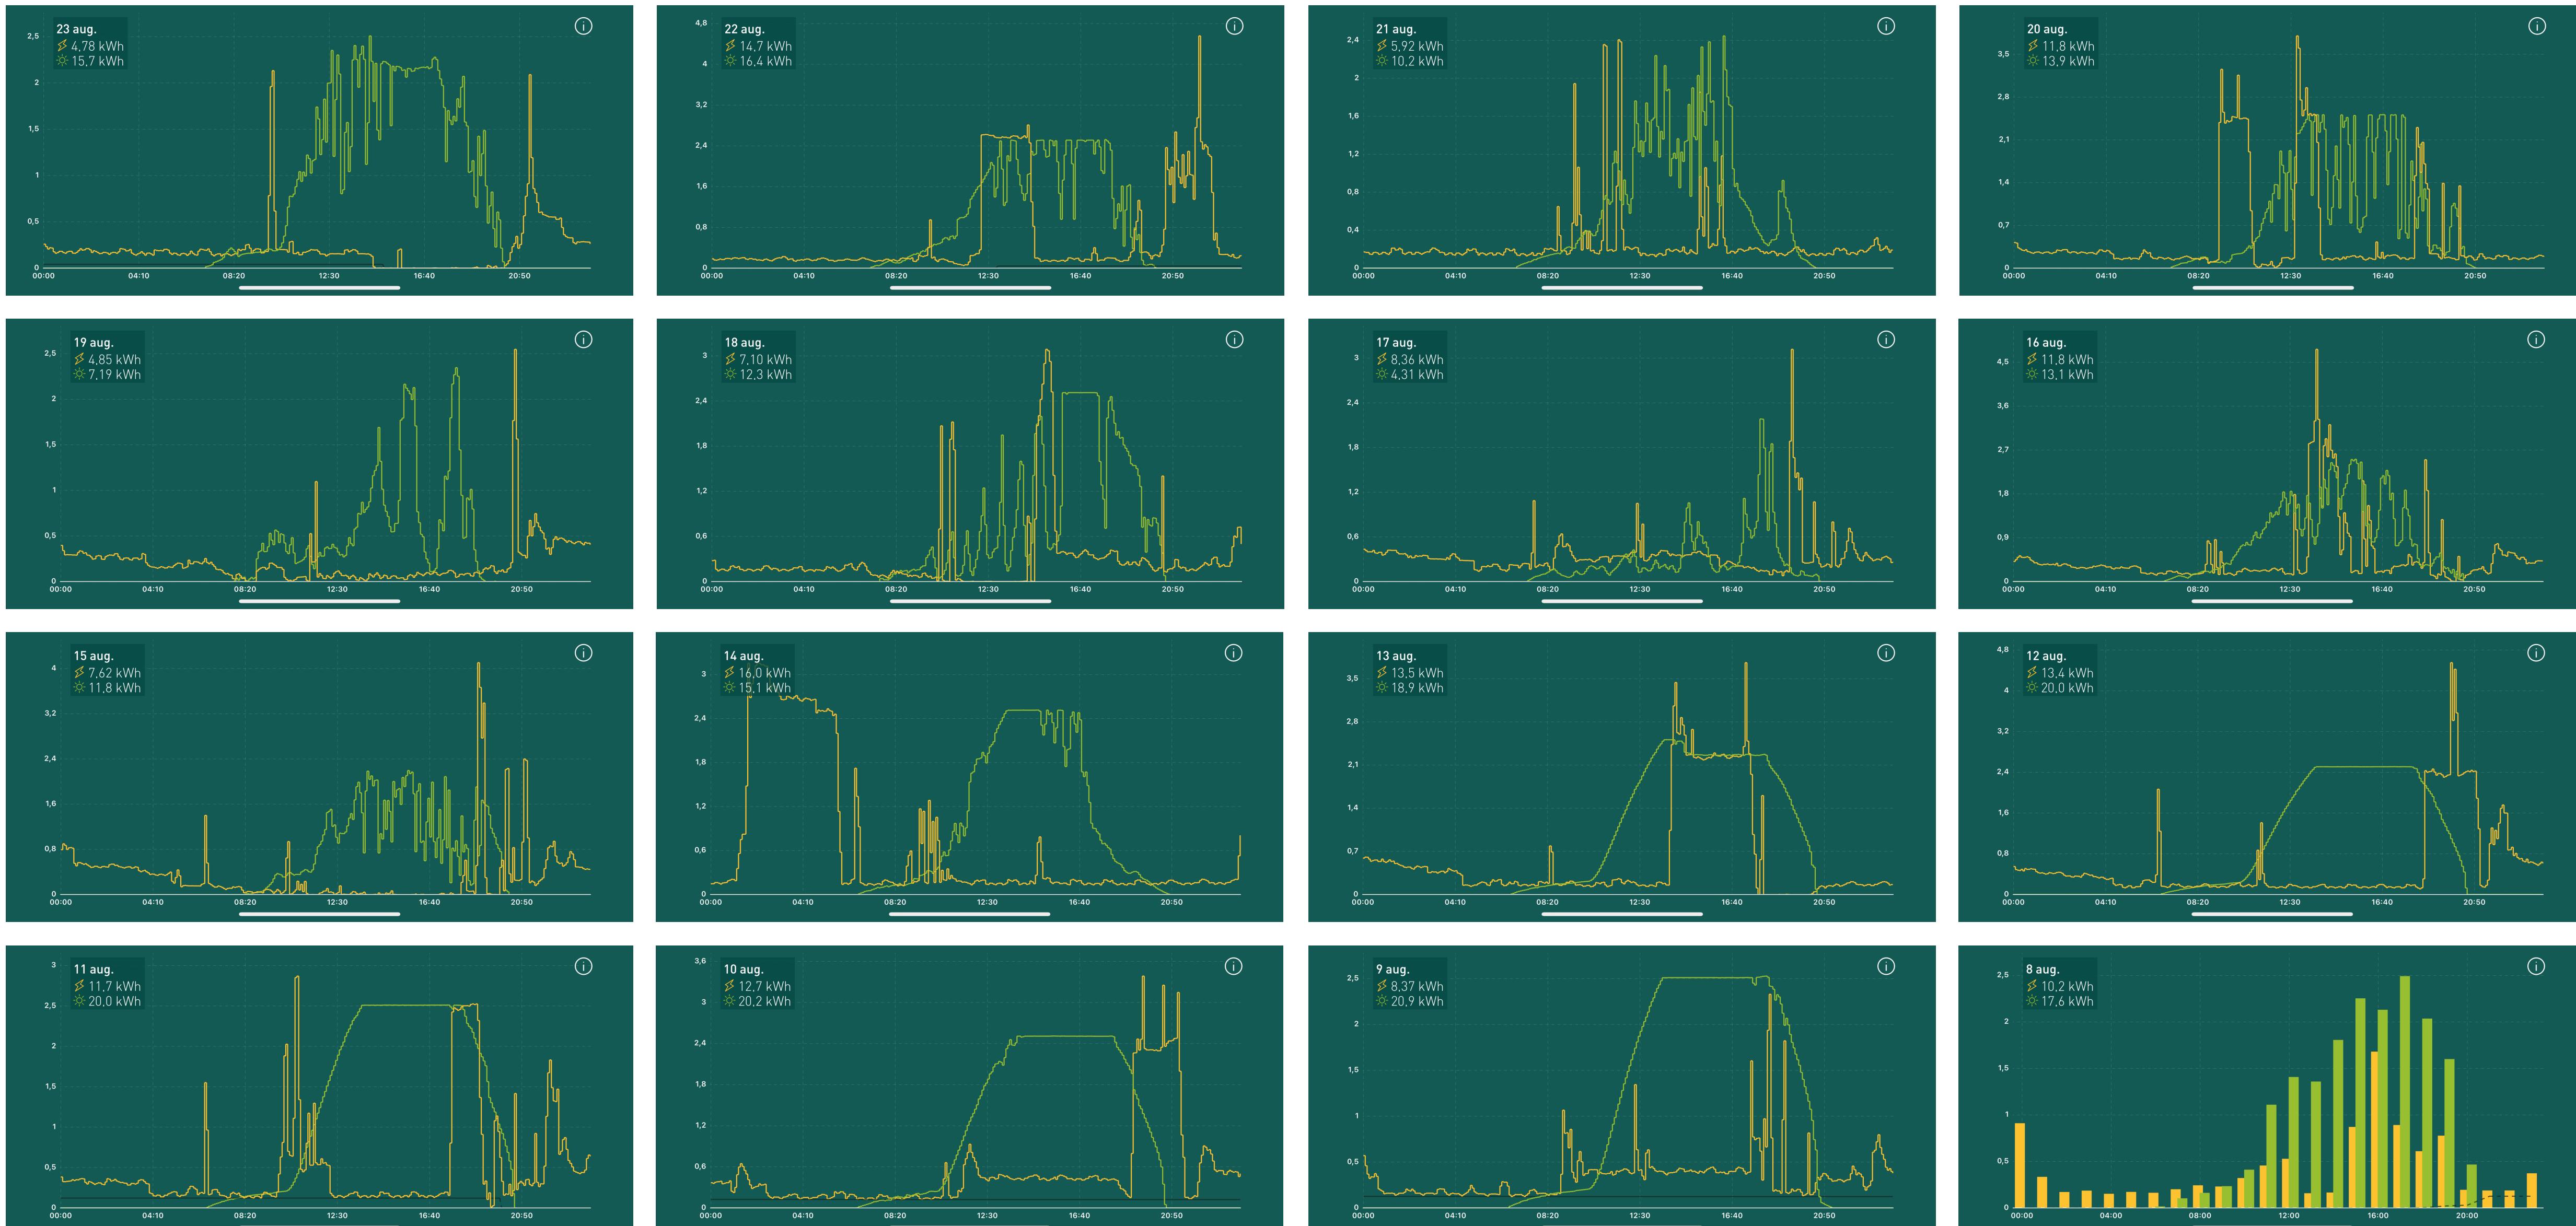
The why...

- Smappee keeps historical data for some time.
- Finest granularity are 5 minute averages:

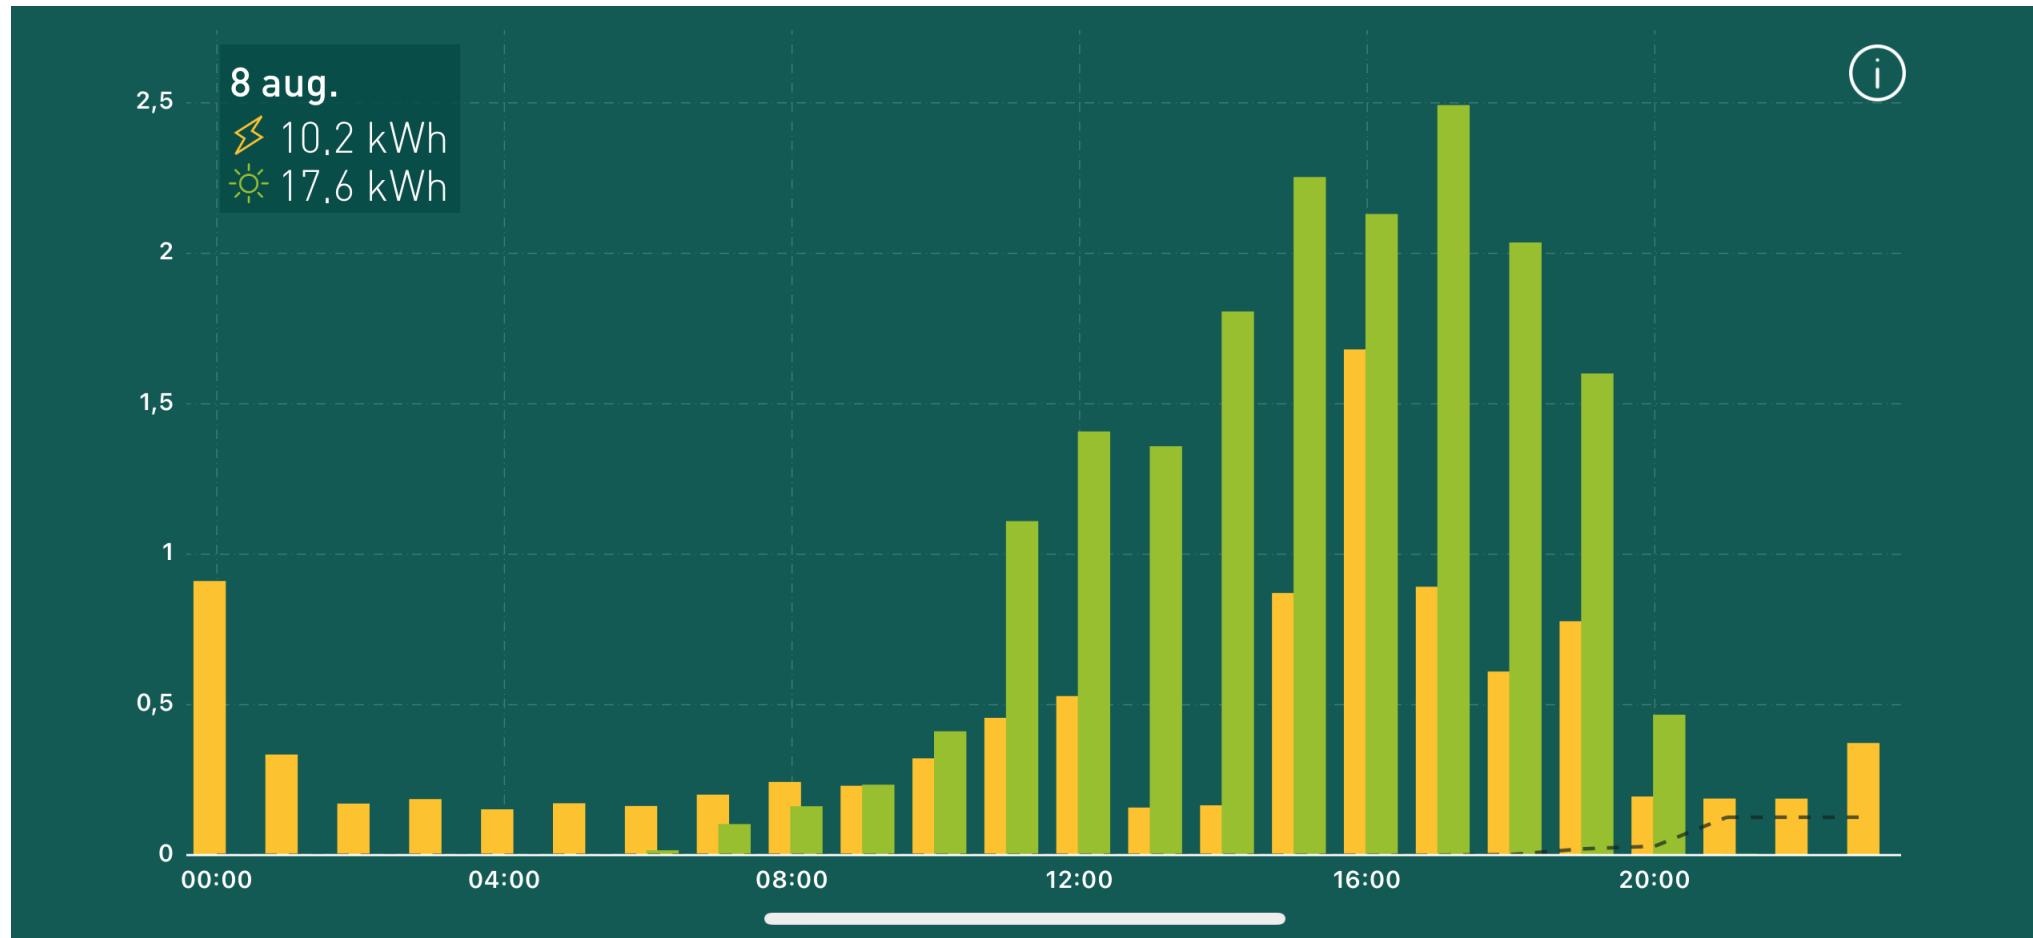


The why...

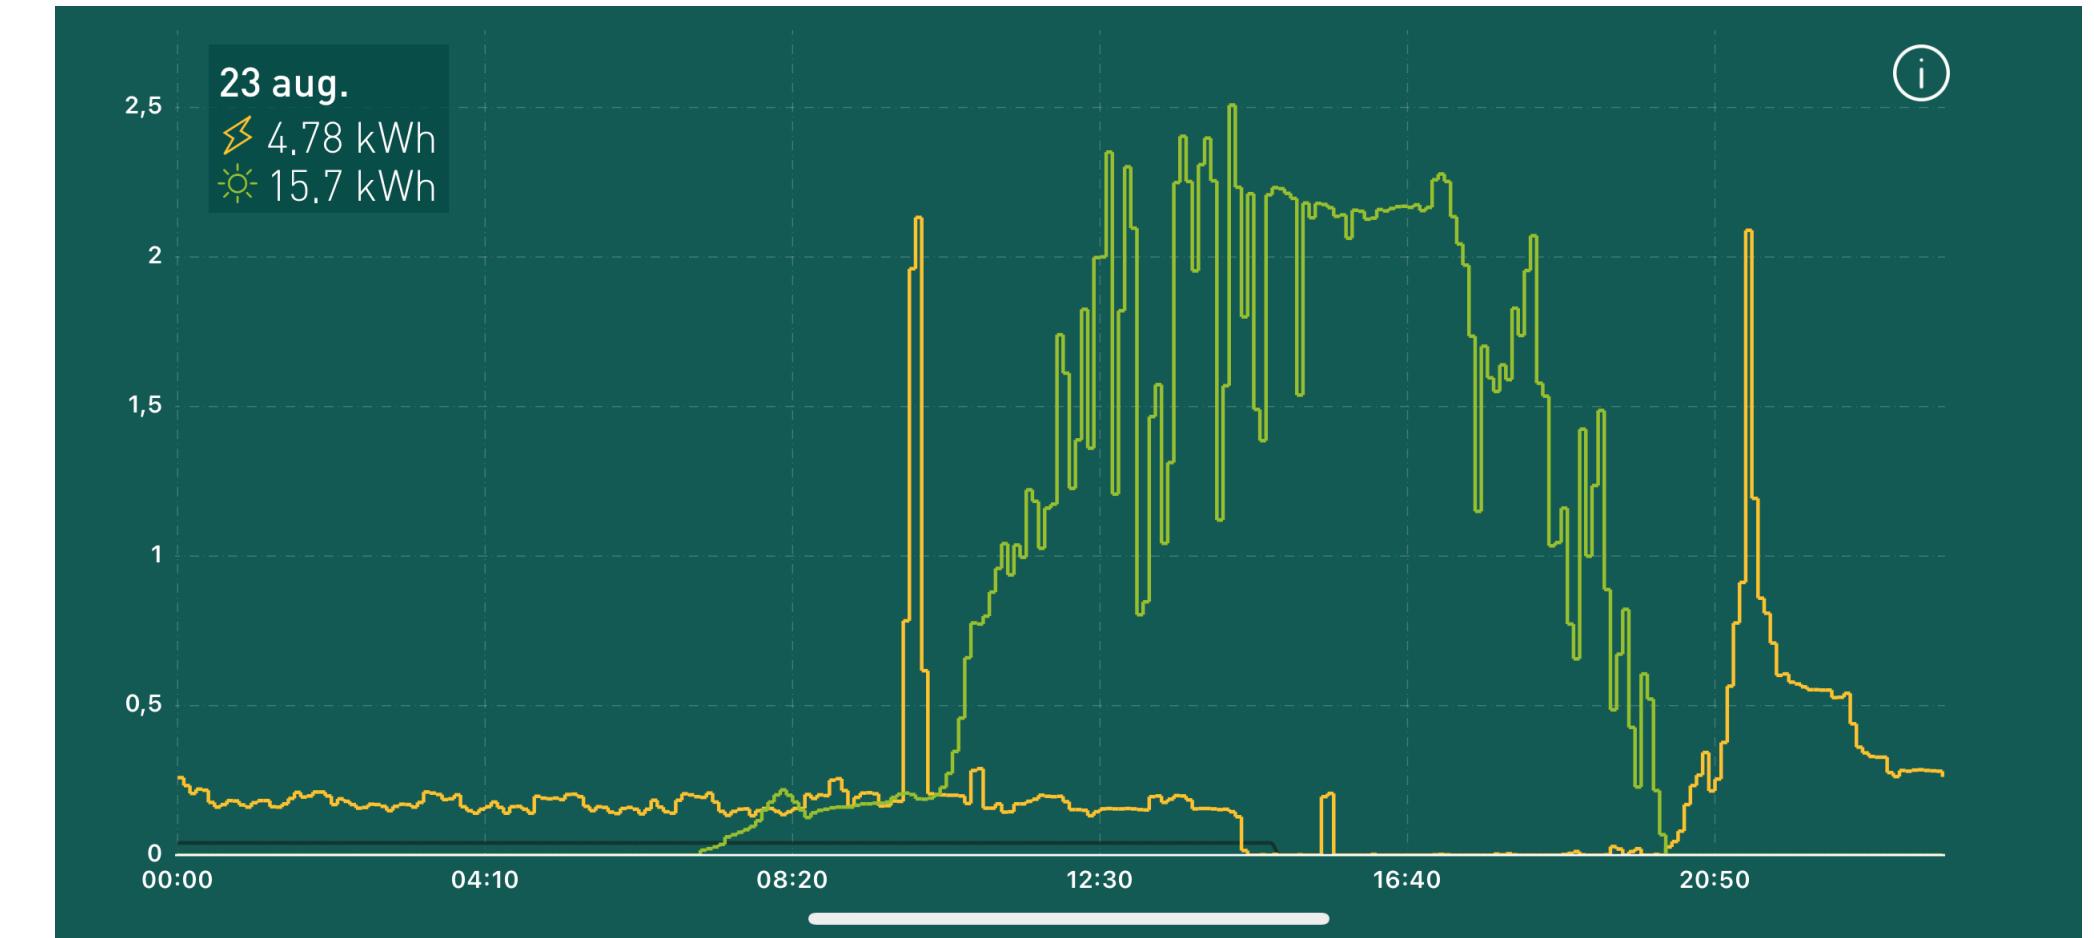
The why...



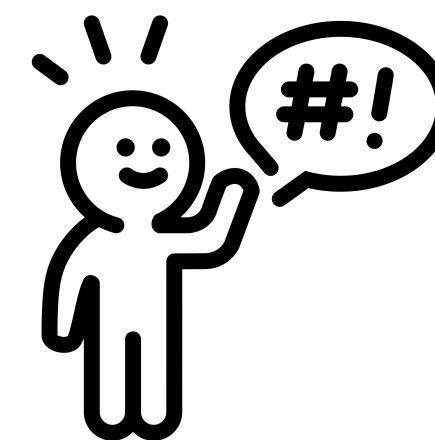
The why...



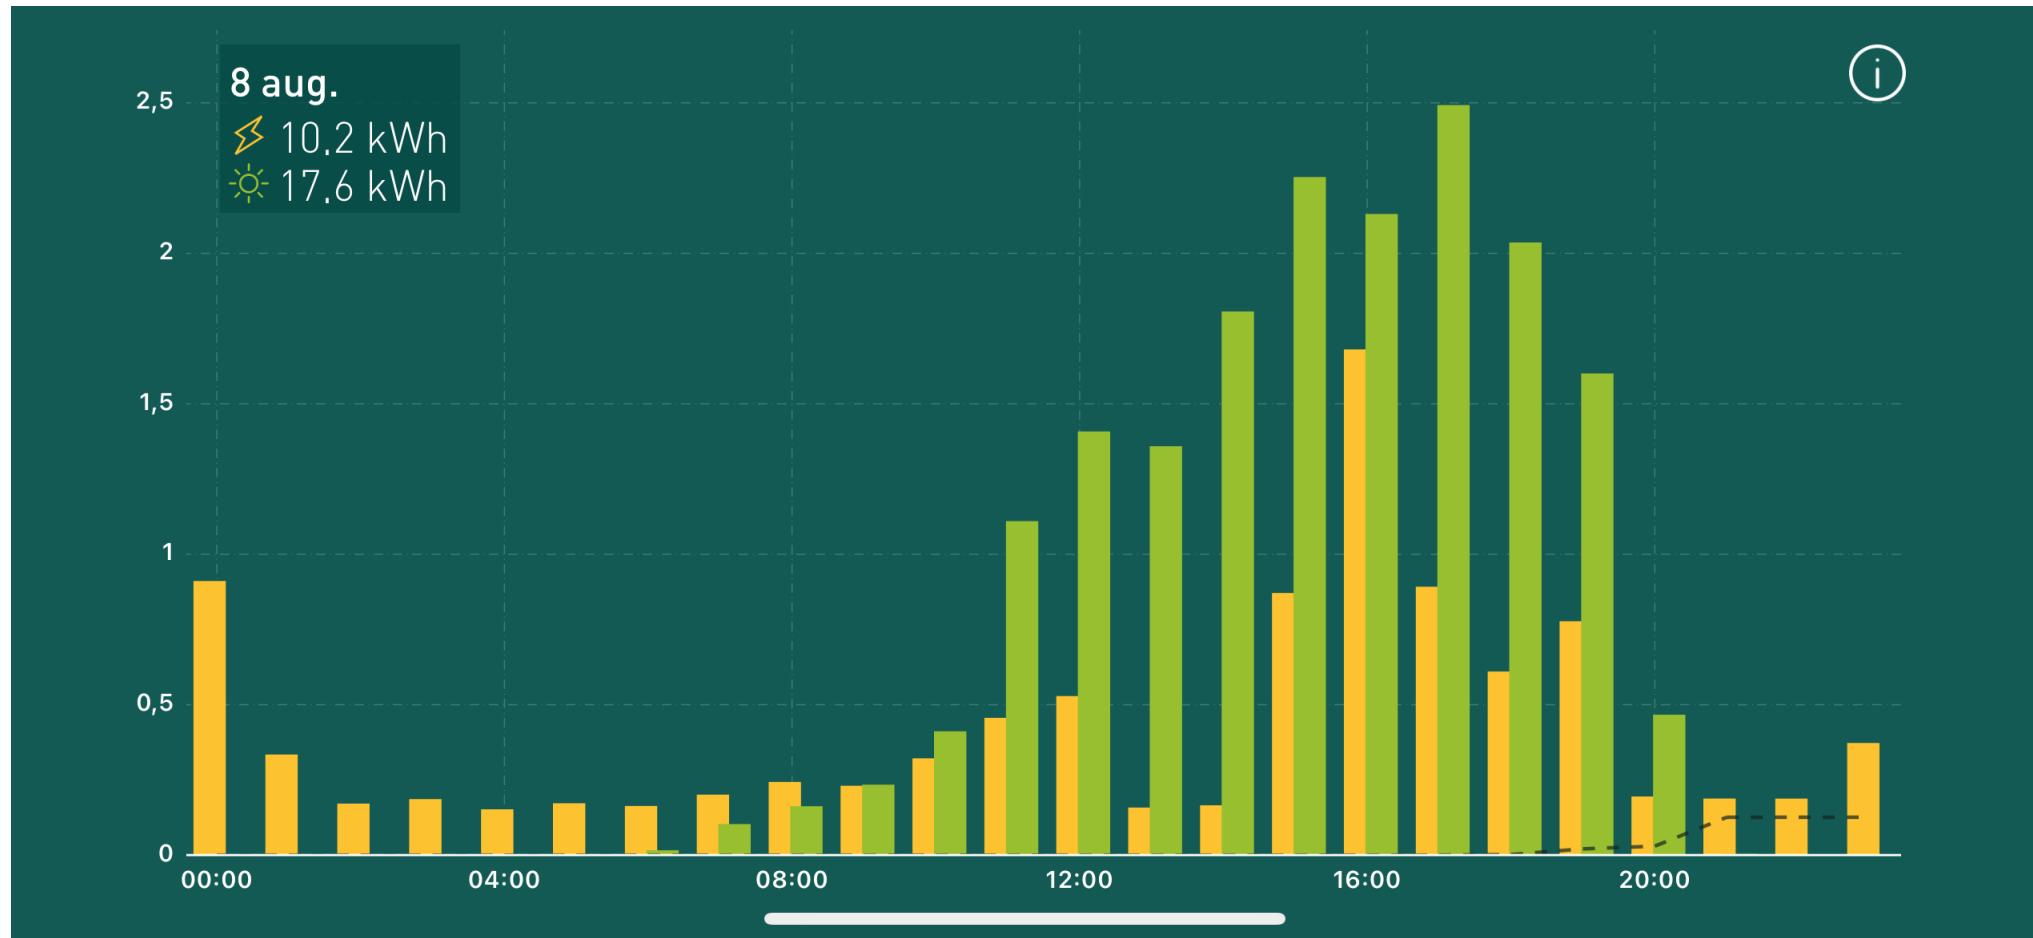
↙ ± 2 weeks



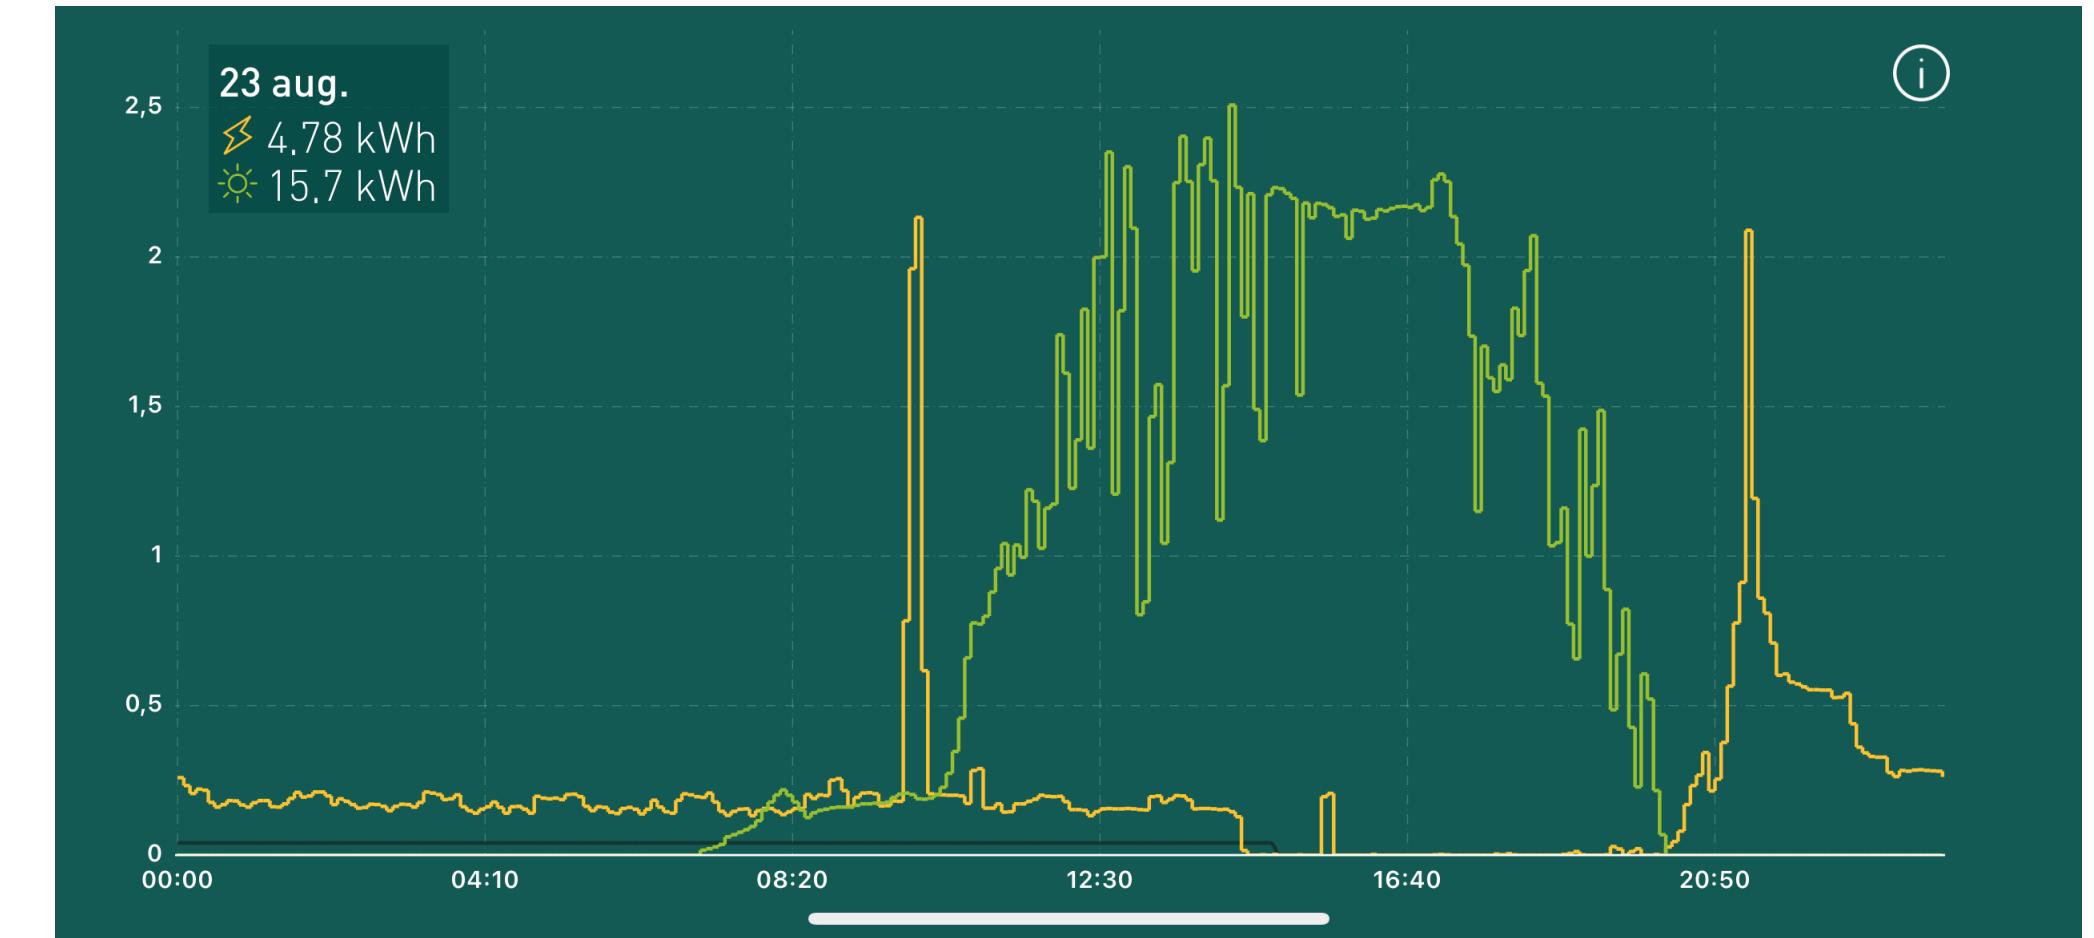
After ± two weeks: 1 hour averages instead of 5 minute averages !



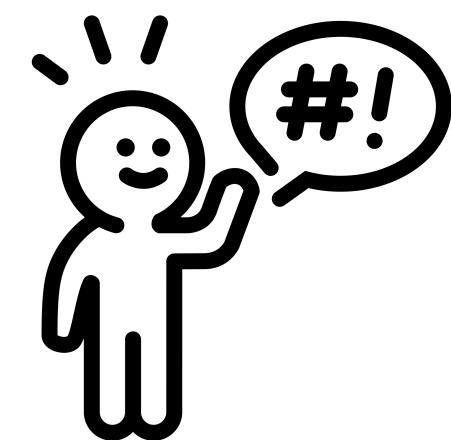
The why...



↙ ± 2 weeks



After ± two weeks: 1 hour averages instead of 5 minute averages !



The why...

what shall we do about this ?

Complain to Smappee ?

Roll our own ?

The why...

what shall we do about this ?

Complain to Smappee ?

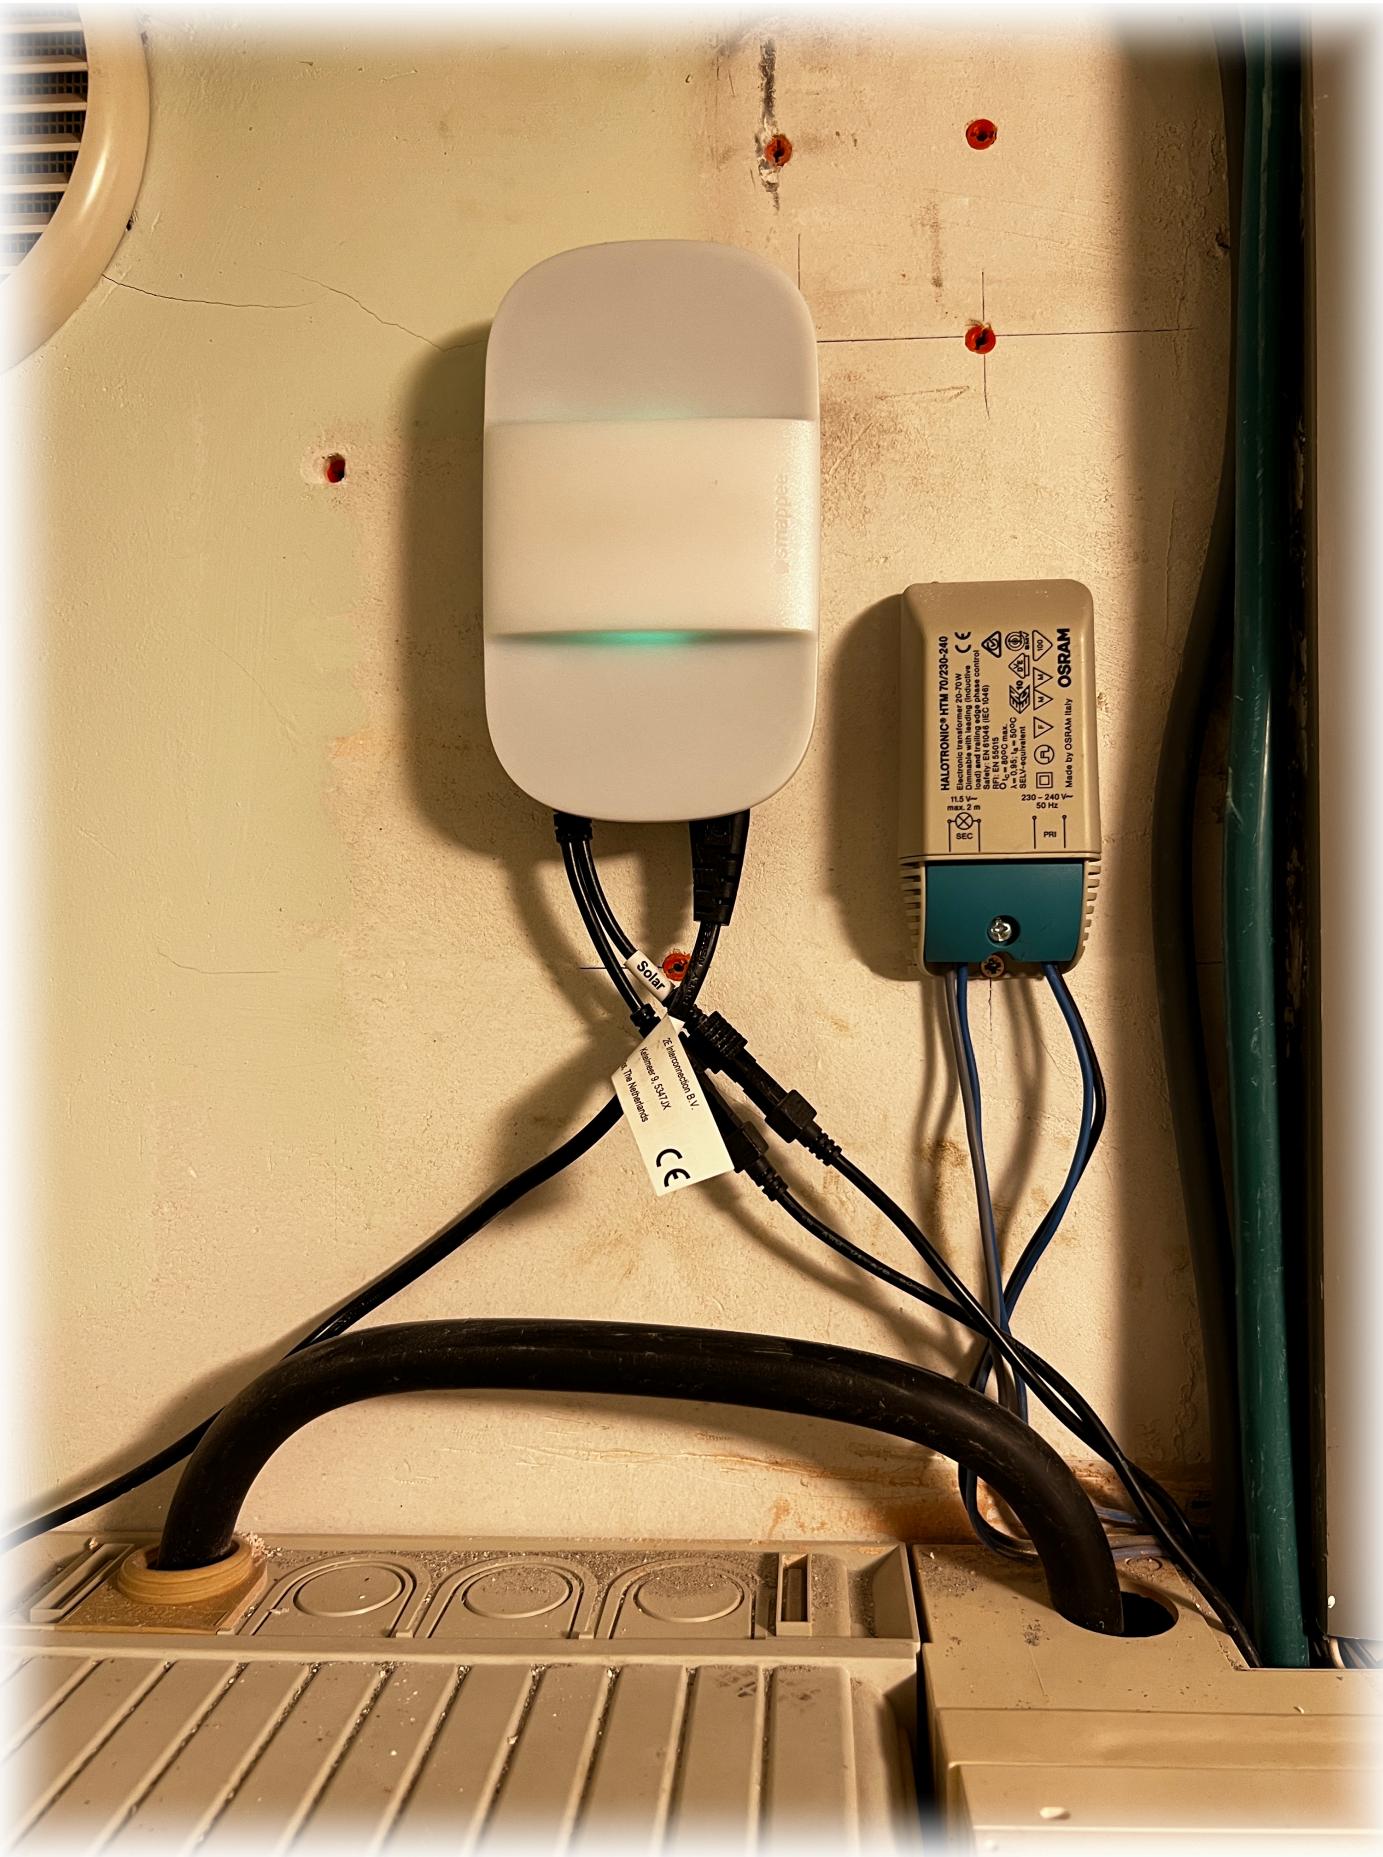
Roll our own ?

Let's explore rolling our own !

Goals of the project

- Run all the bits on a Raspberry Pi with an USB-attached SSD
- Implement the monitoring system that
 - avoids programming if possible
 - uses Open Source Software only
 - captures all knowledge as to make this easily reproducible [for others]
 - automates the installation of the OS and required software

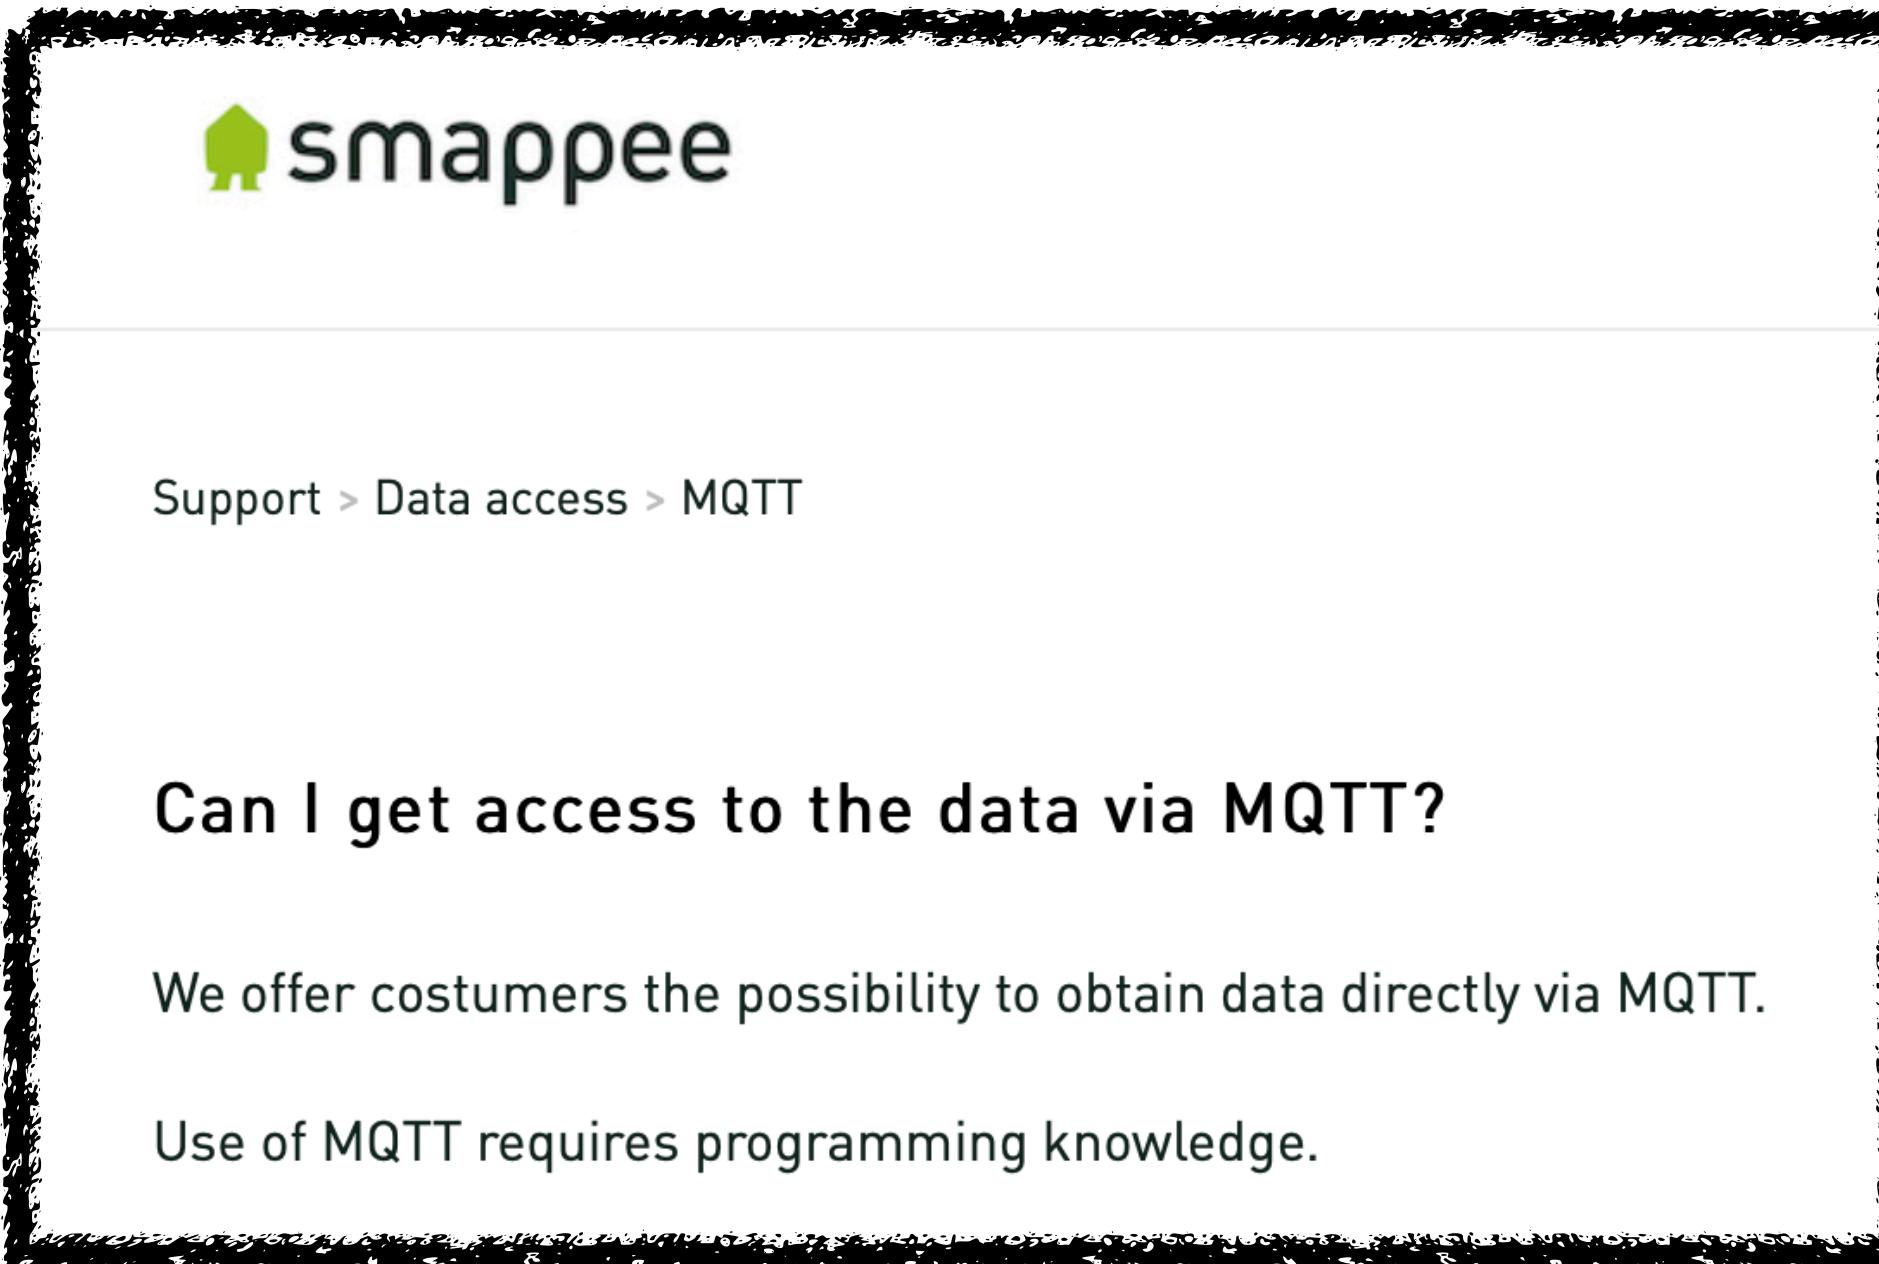
The Smappee device



- 4-channel power measurement
- 1 second sample interval
- MQTT client

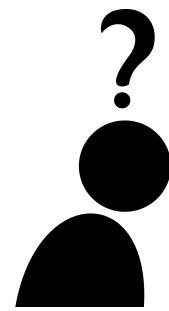
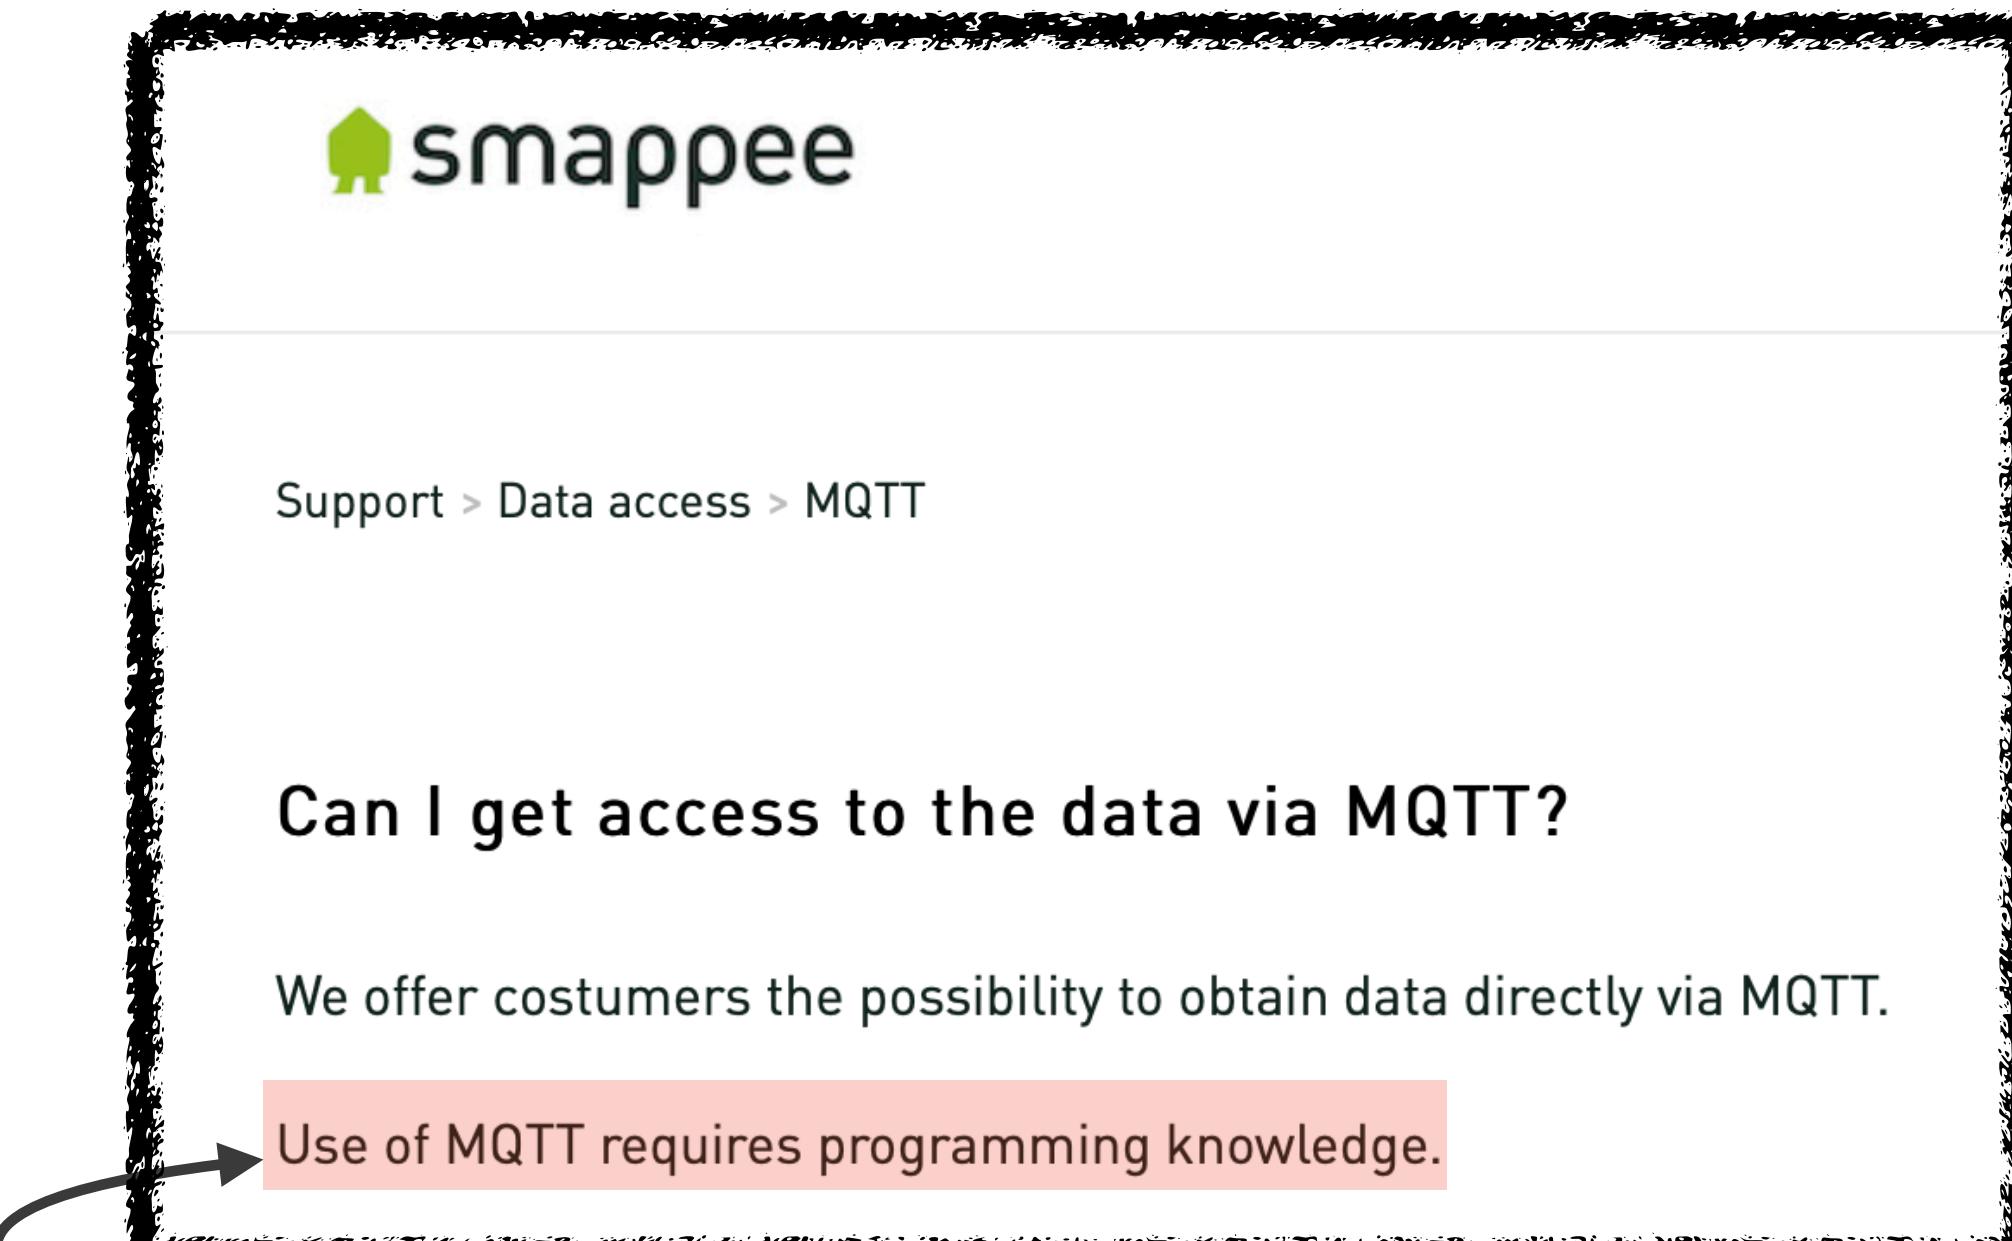
What are our options?

- Smappee offers access to data in two ways
 - via REST API (paying service)
 - via MQTT:



What are our options?

- Smappee offers access to data in two ways
 - via REST API (paying service)
 - via MQTT:

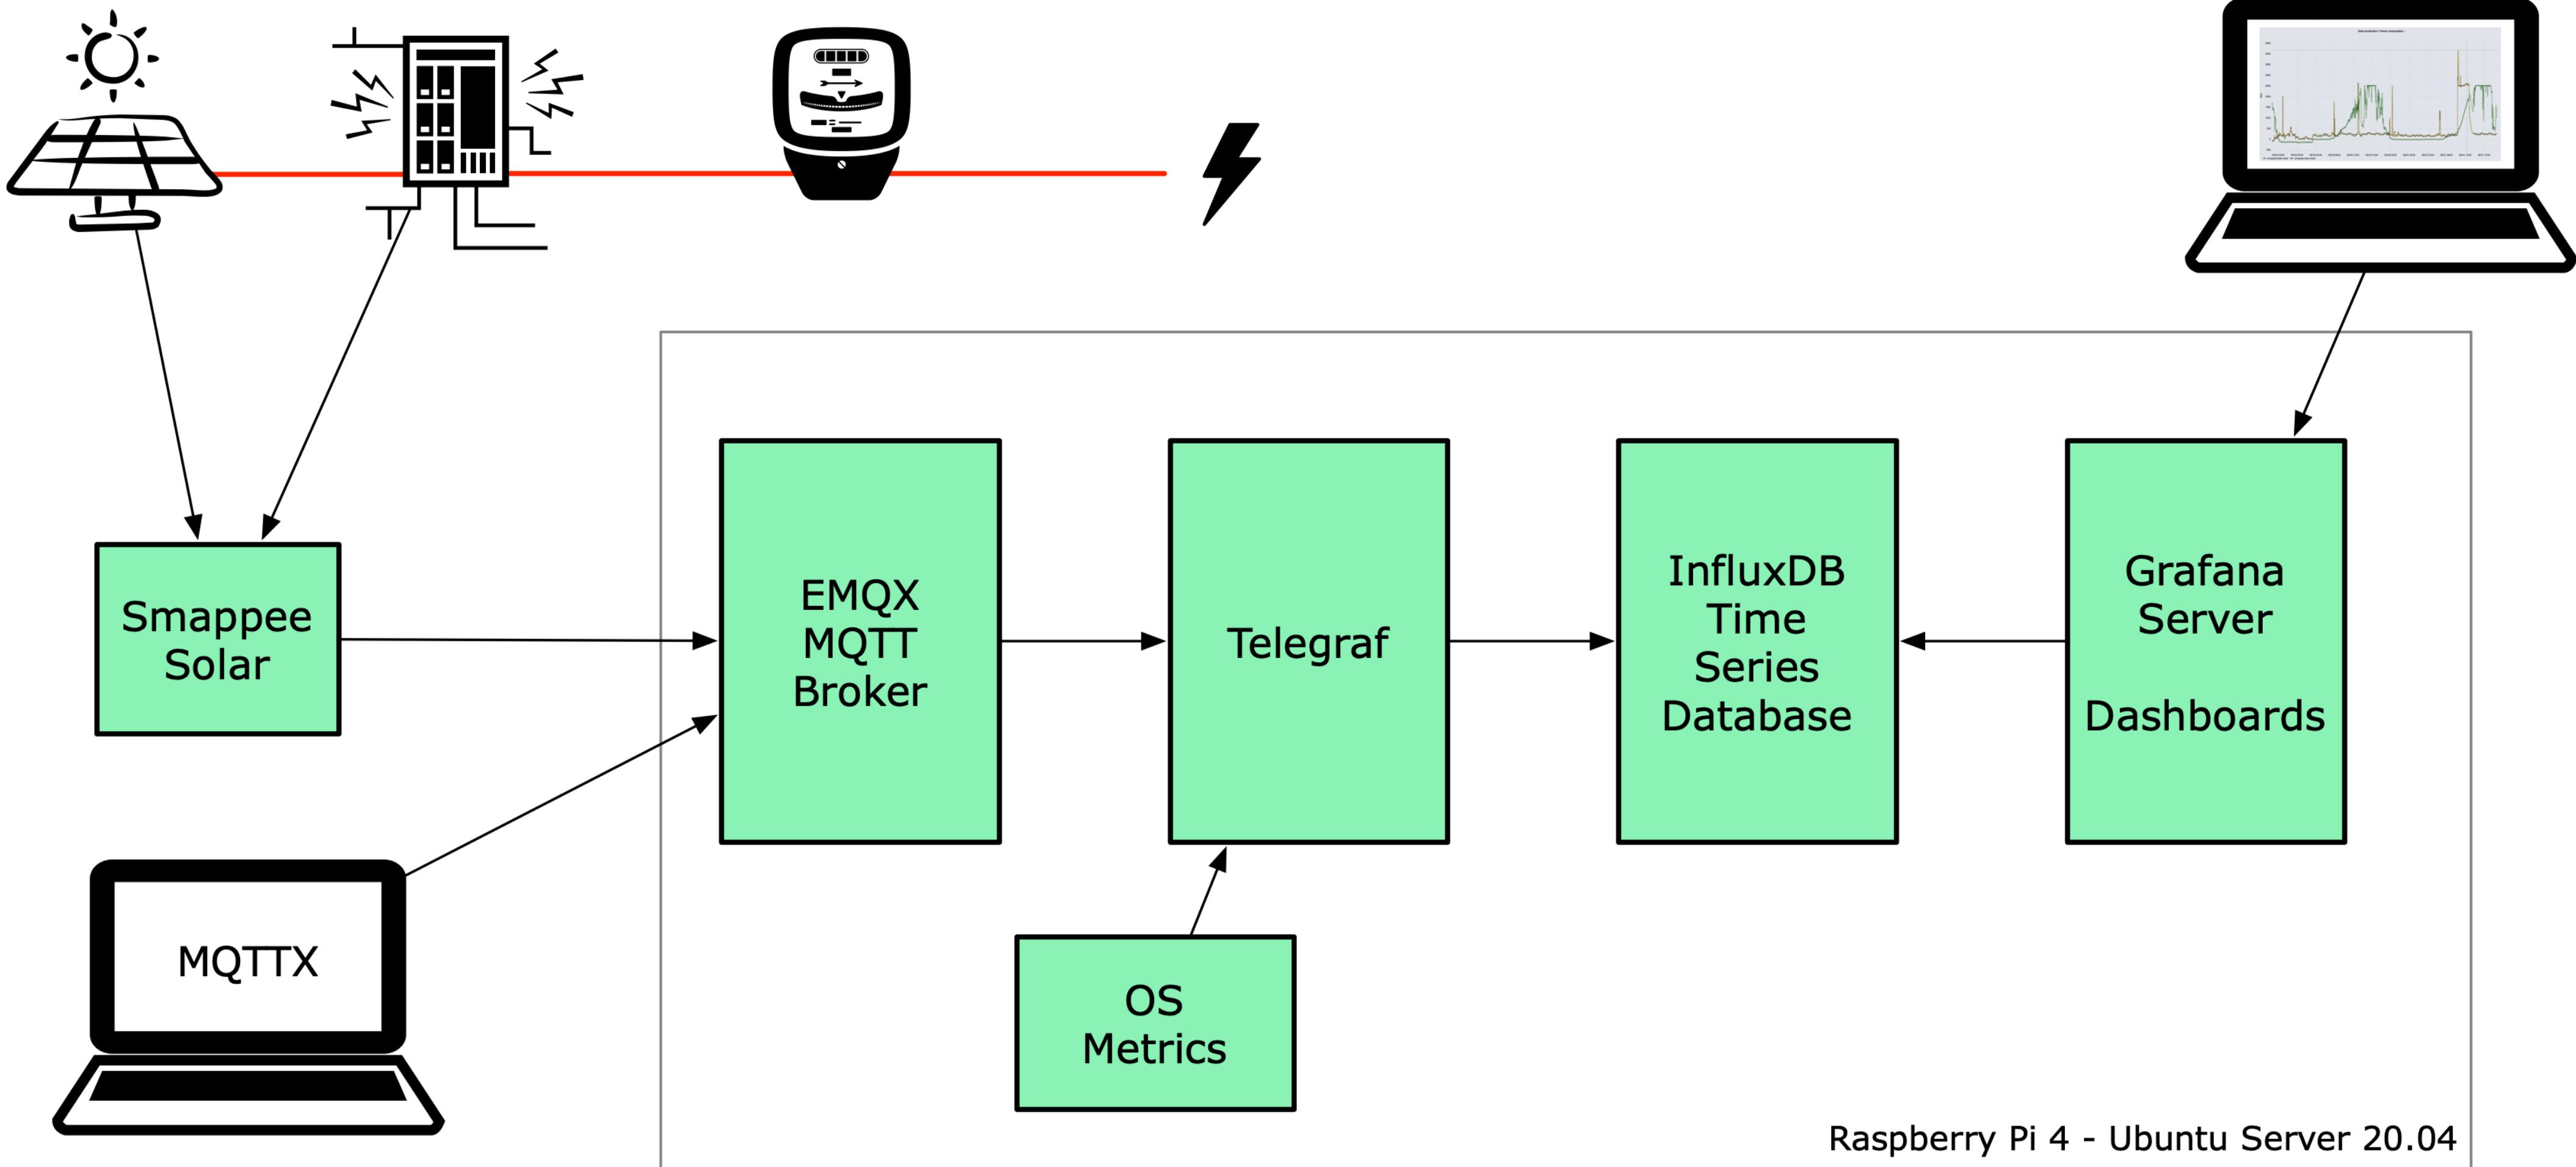


Sounds interesting, but do I really need to program?

The base platform

- Hardware:
 - Raspberry Pi 4 Model B - 8GB (smaller will do)
 - Storage: Samsung 250GB SSD attached to a USB-3 port
- OS:
 - 64-bit Ubuntu Server 20.04.4 LTS
 - Has cloud-init (version 21.4) support
 - Can boot off a USB disk with a gotcha

Home monitoring system setup



The software components

Component	Description	Version	OS support
EMQX	MQTT broker	5.0.4	Ubuntu 20.04
MQTTX	Cross-platform MQTT 5.0 client tool, GUI & CLI	1.6.0	*nix, Mac, Windows
InfluxDB	Time series database	1.8.10	≥ Ubuntu 20.04
Telegraf	An agent for collecting, processing, aggregating, and writing metrics	1.23.3	≥ Ubuntu 20.04
Grafana Server	Query, visualize, alert on and understand your metrics	9.0.6	

EMQX forces the utilisation of Ubuntu 20.04

Configuring/testing MQTT

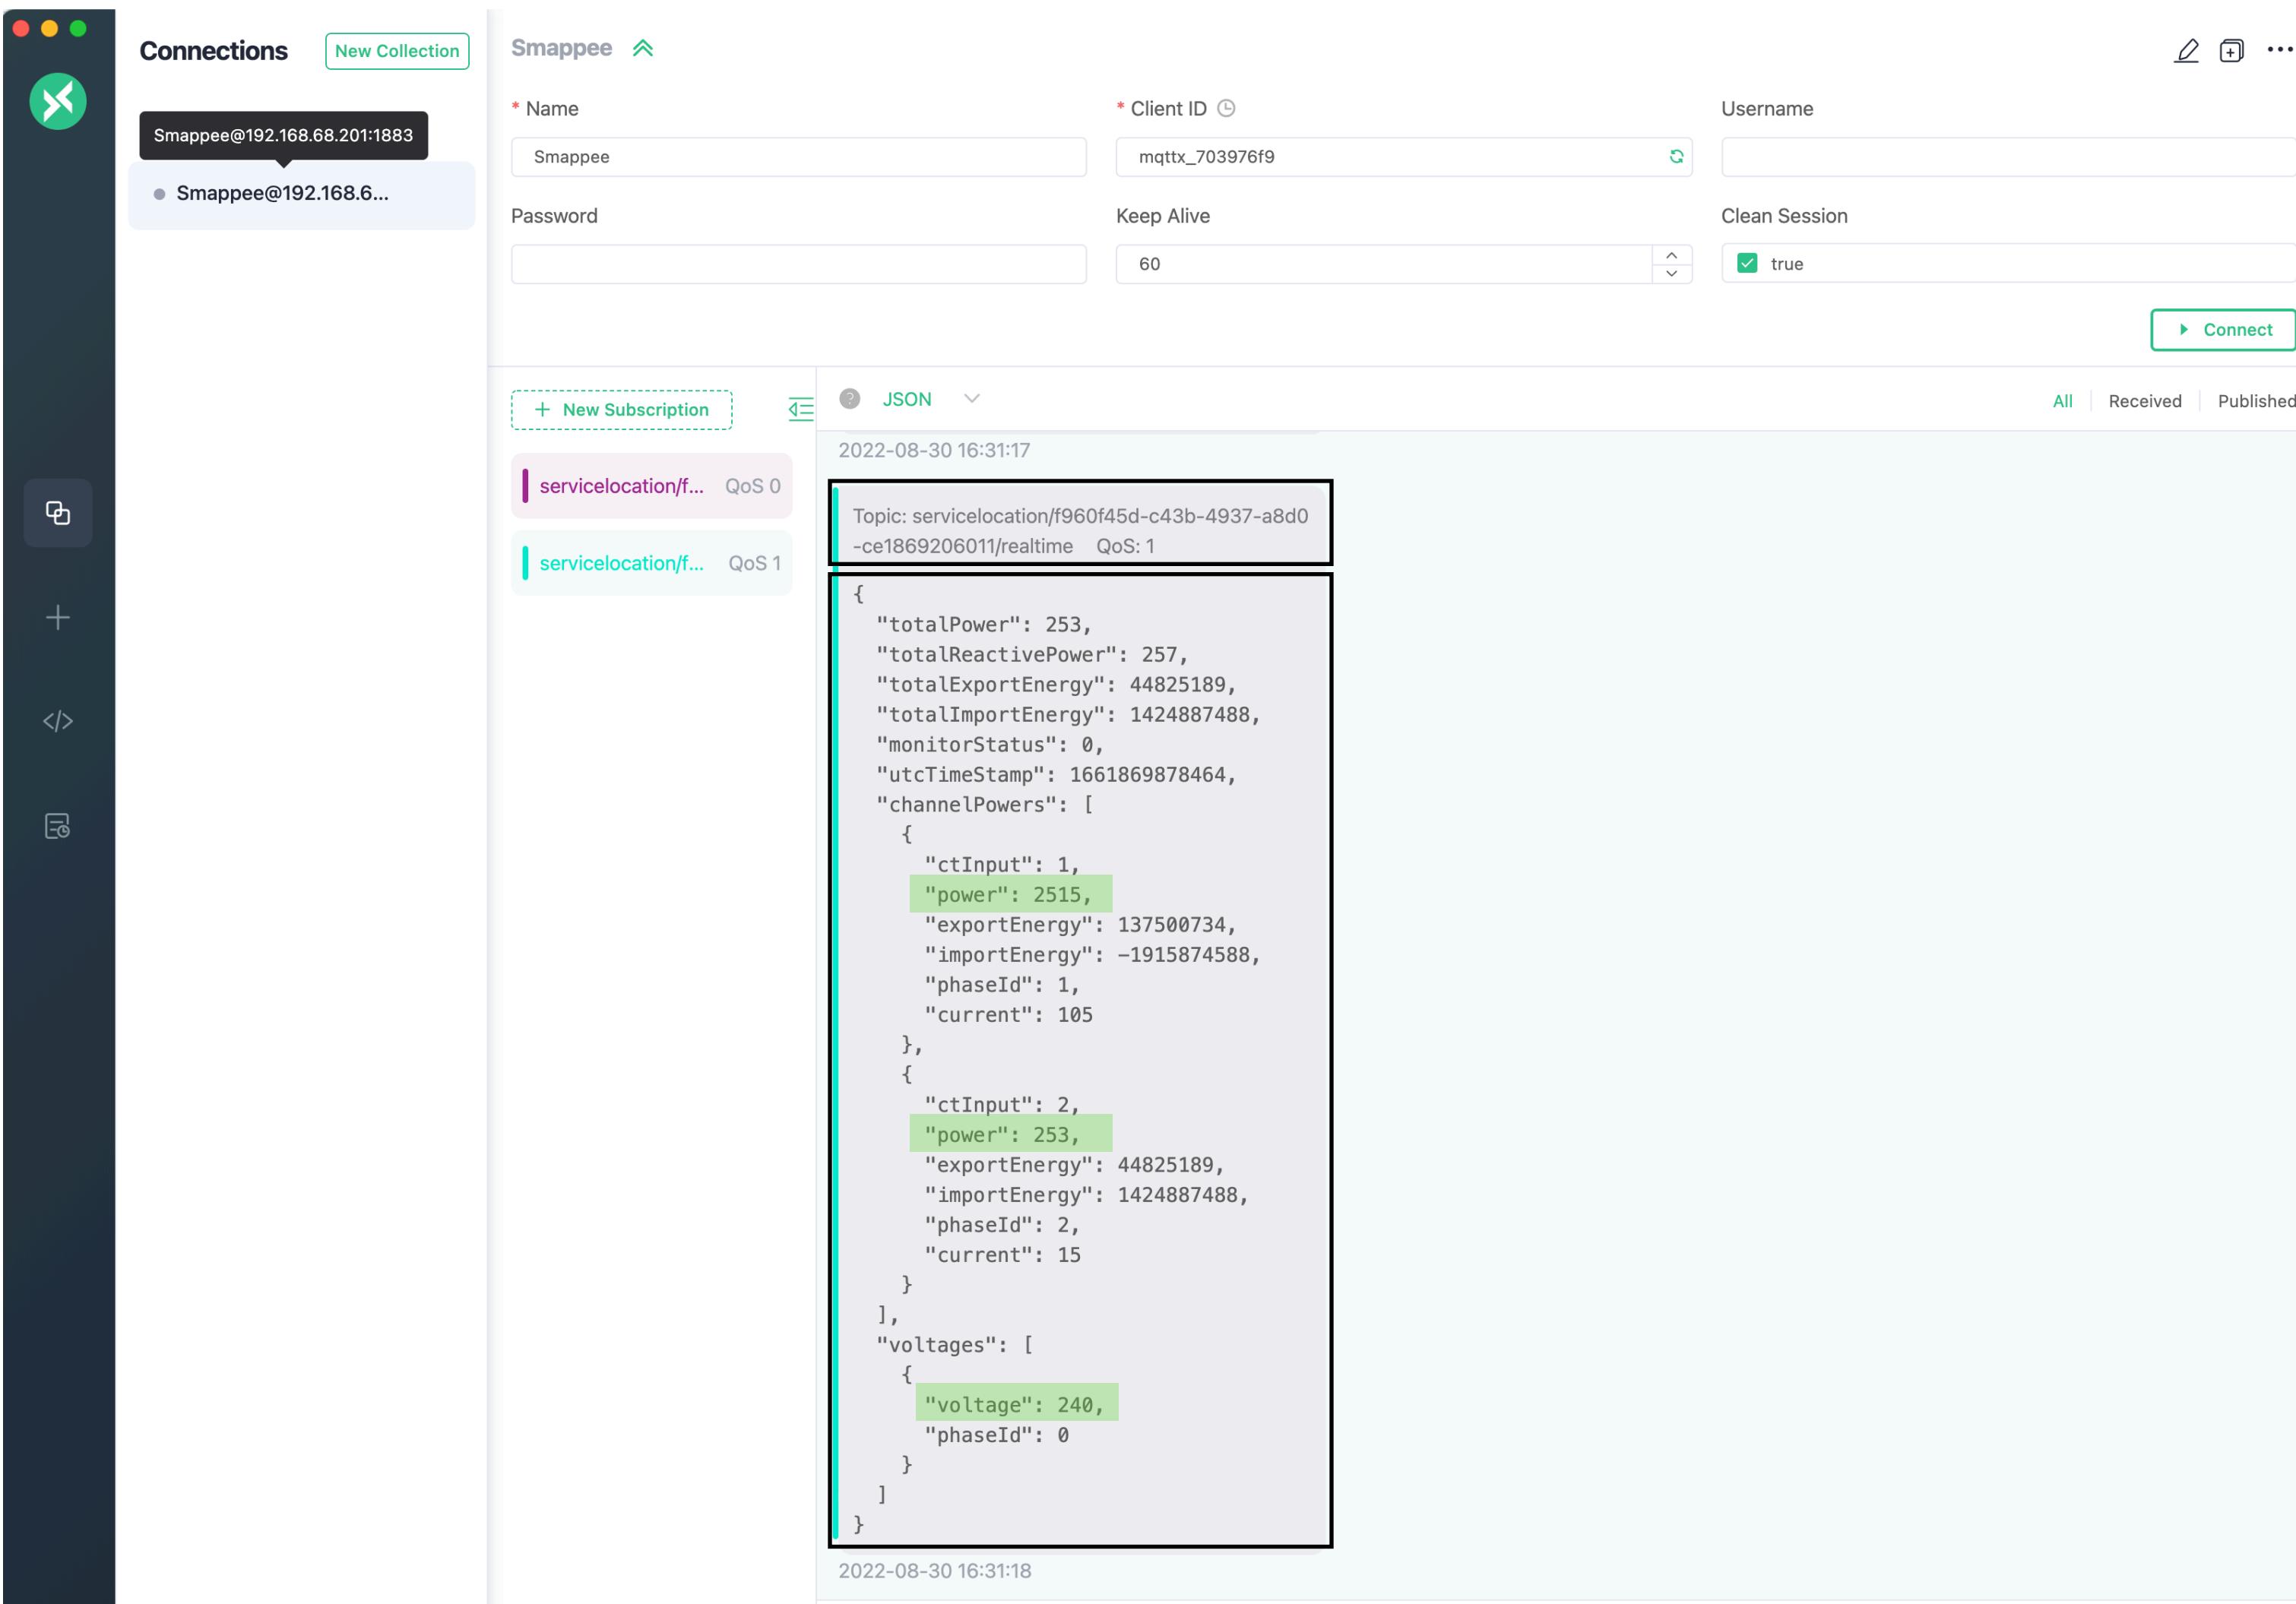
- We configure Smappee for our MQTT Broker (EMQX)
our broker lives at 192.168.68.201:1883

Advanced config parameters	Setting
Webportal password	admin
Active power lower limit (color green->yellow)	0
Active power upper limit (color yellow->orange)	0
MDNS hostname	Smappee11070075
MQTT local broker (e.g. tcp://192.168.0.48:1883)	tcp://192.168.68.20
MQTT local broker username	
MQTT local broker password	
Switch 5 minute read delay	0
<input type="button" value="Apply changes and restart monitor"/>	

Advanced action	
Restart monitor	<input type="button" value="apply"/>
Clear all data (historical data + appliances data) and restart monitor	<input type="button" value="apply"/>
Clear only appliances data and restart monitor	<input type="button" value="apply"/>
Reset active power peak values	<input type="button" value="apply"/>
Reset sensor cache	<input type="button" value="apply"/>
Autocommissioning validation	<input type="button" value="stop"/> <input type="button" value="start"/> <input type="button" value="reset values"/>
Enable/disable green breathing	<input checked="" type="checkbox"/>
Enable/disable command control look ahead schedule recover	<input checked="" type="checkbox"/>
Enable/disable MAC address arp-scan	<input type="checkbox"/>
Reset MAC address scanning cache	<input type="button" value="apply"/>
Enable/disable MDNS (monitor will reboot)	<input checked="" type="checkbox"/>
Advanced	<input checked="" type="checkbox"/>

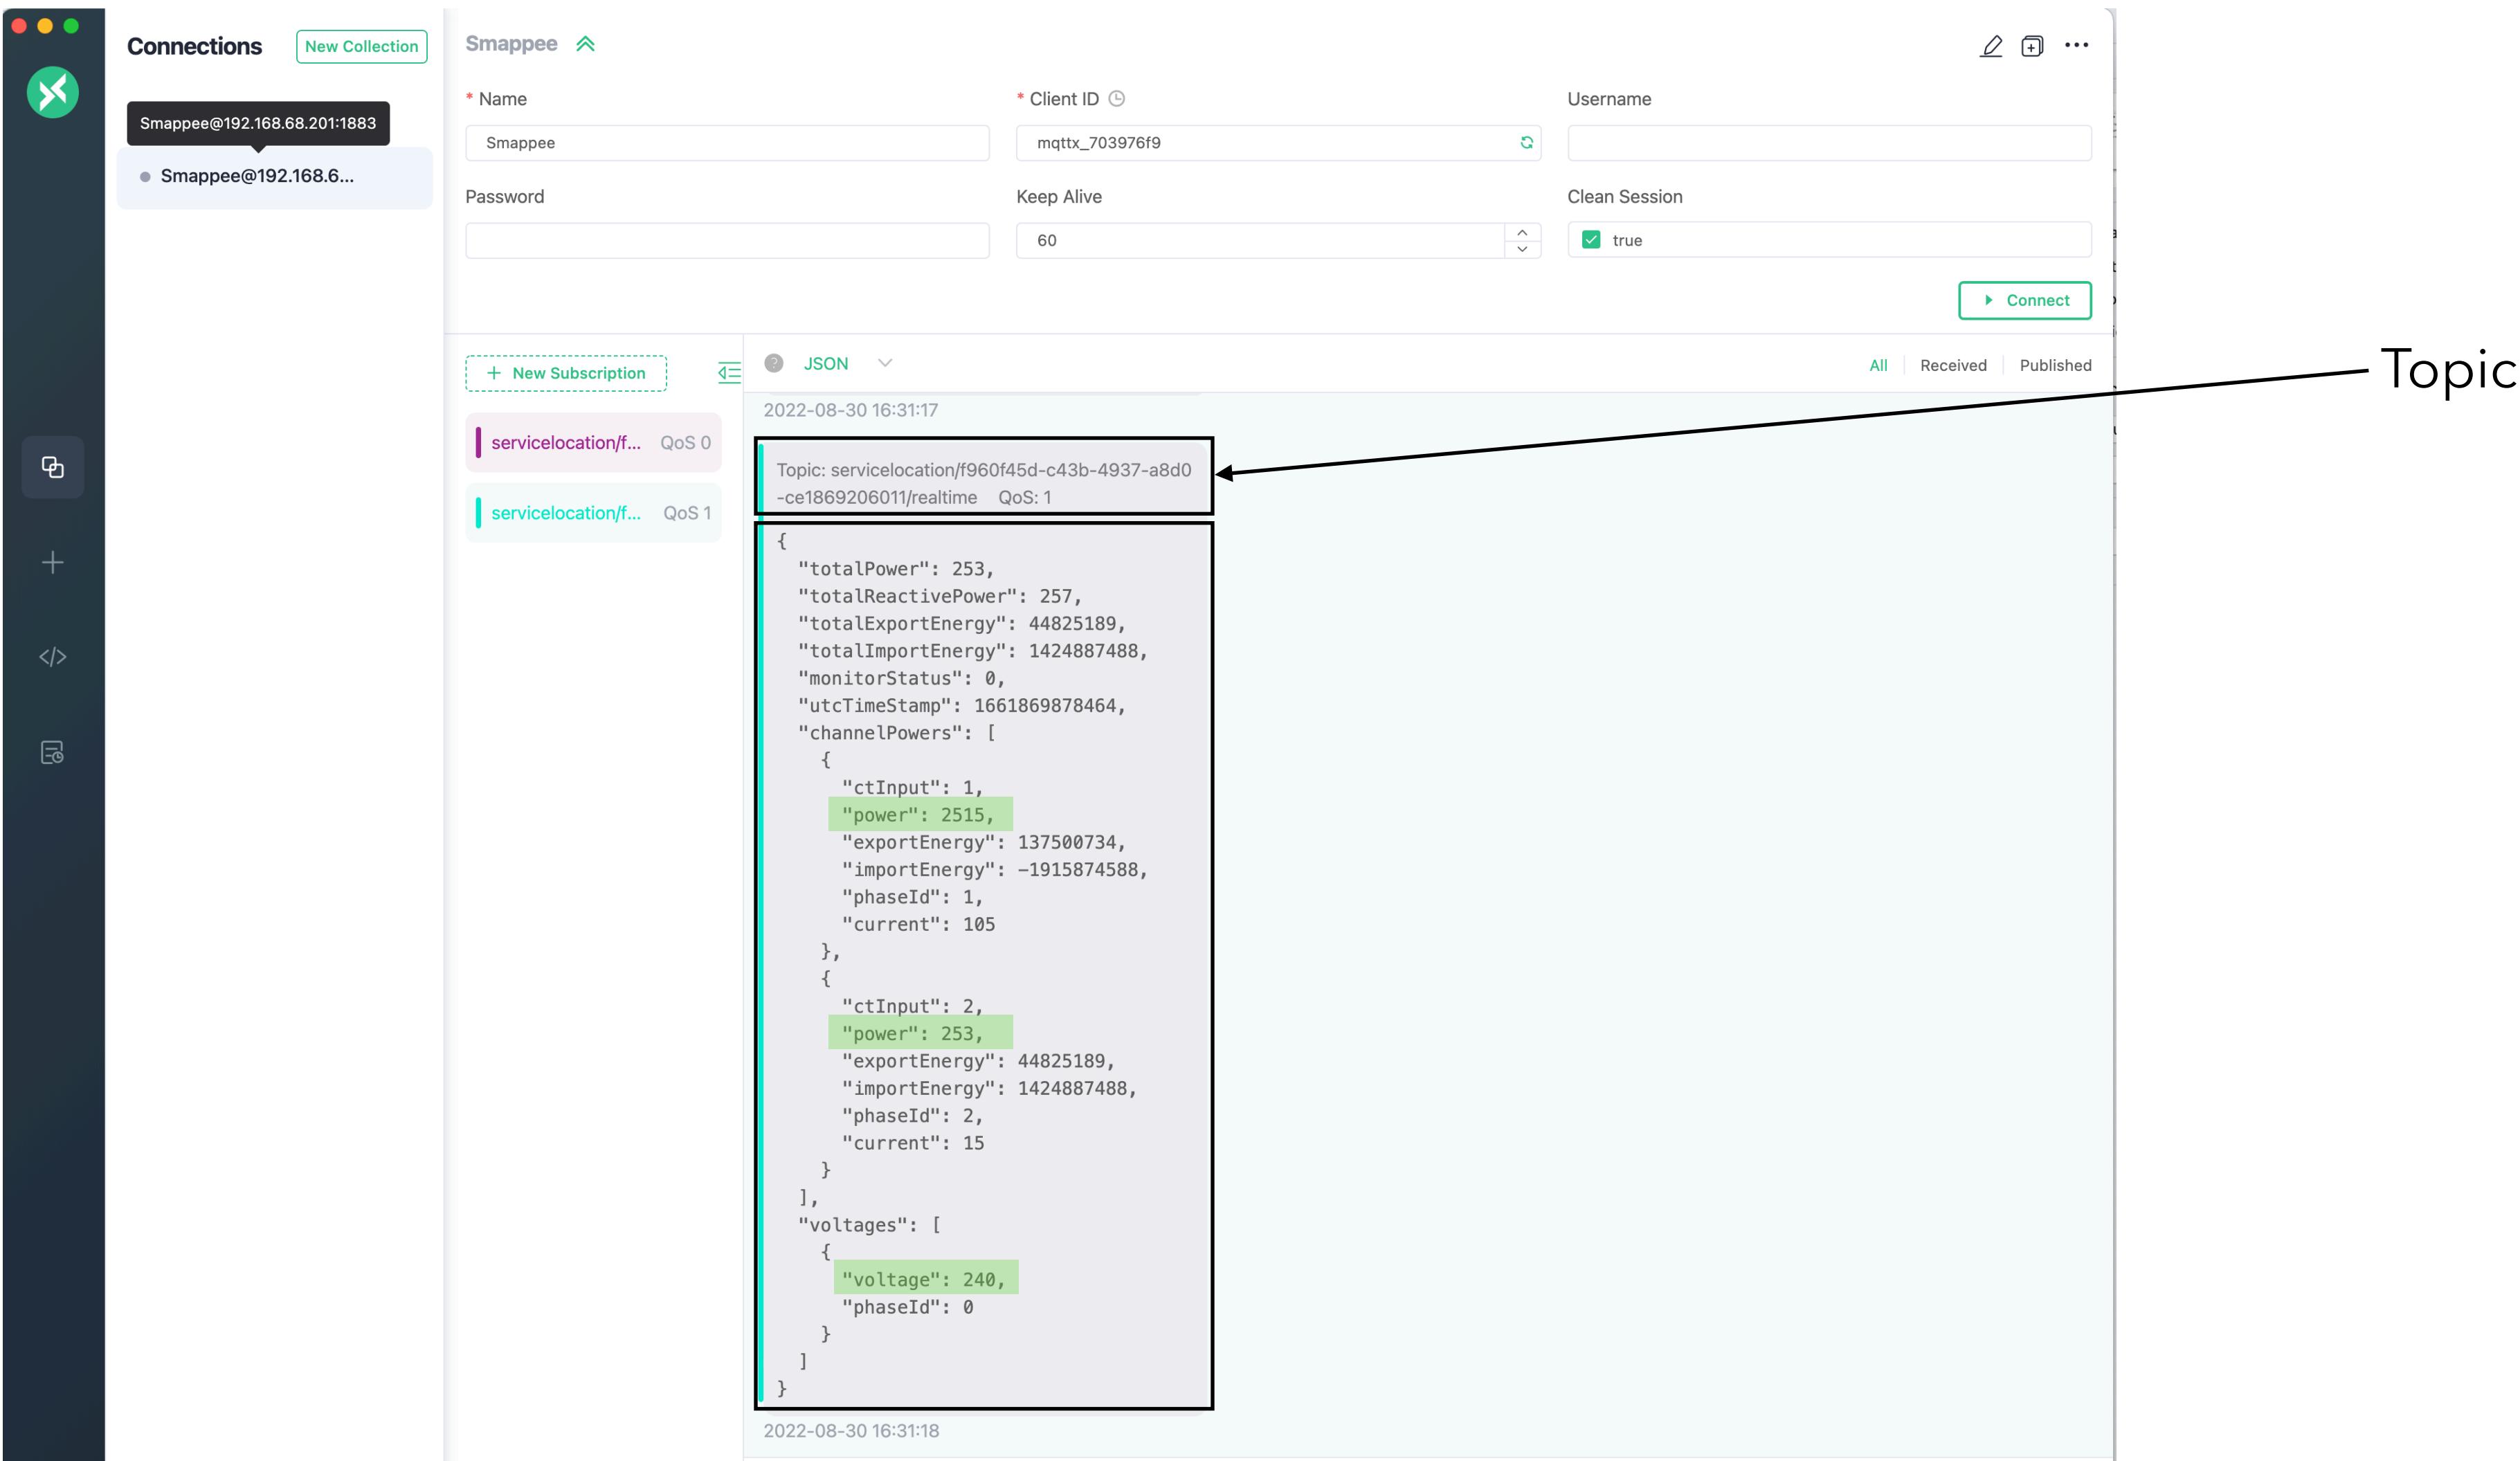
Configuring/testing MQTT

- Let's use MQTTX to find out what data our Smappee sends



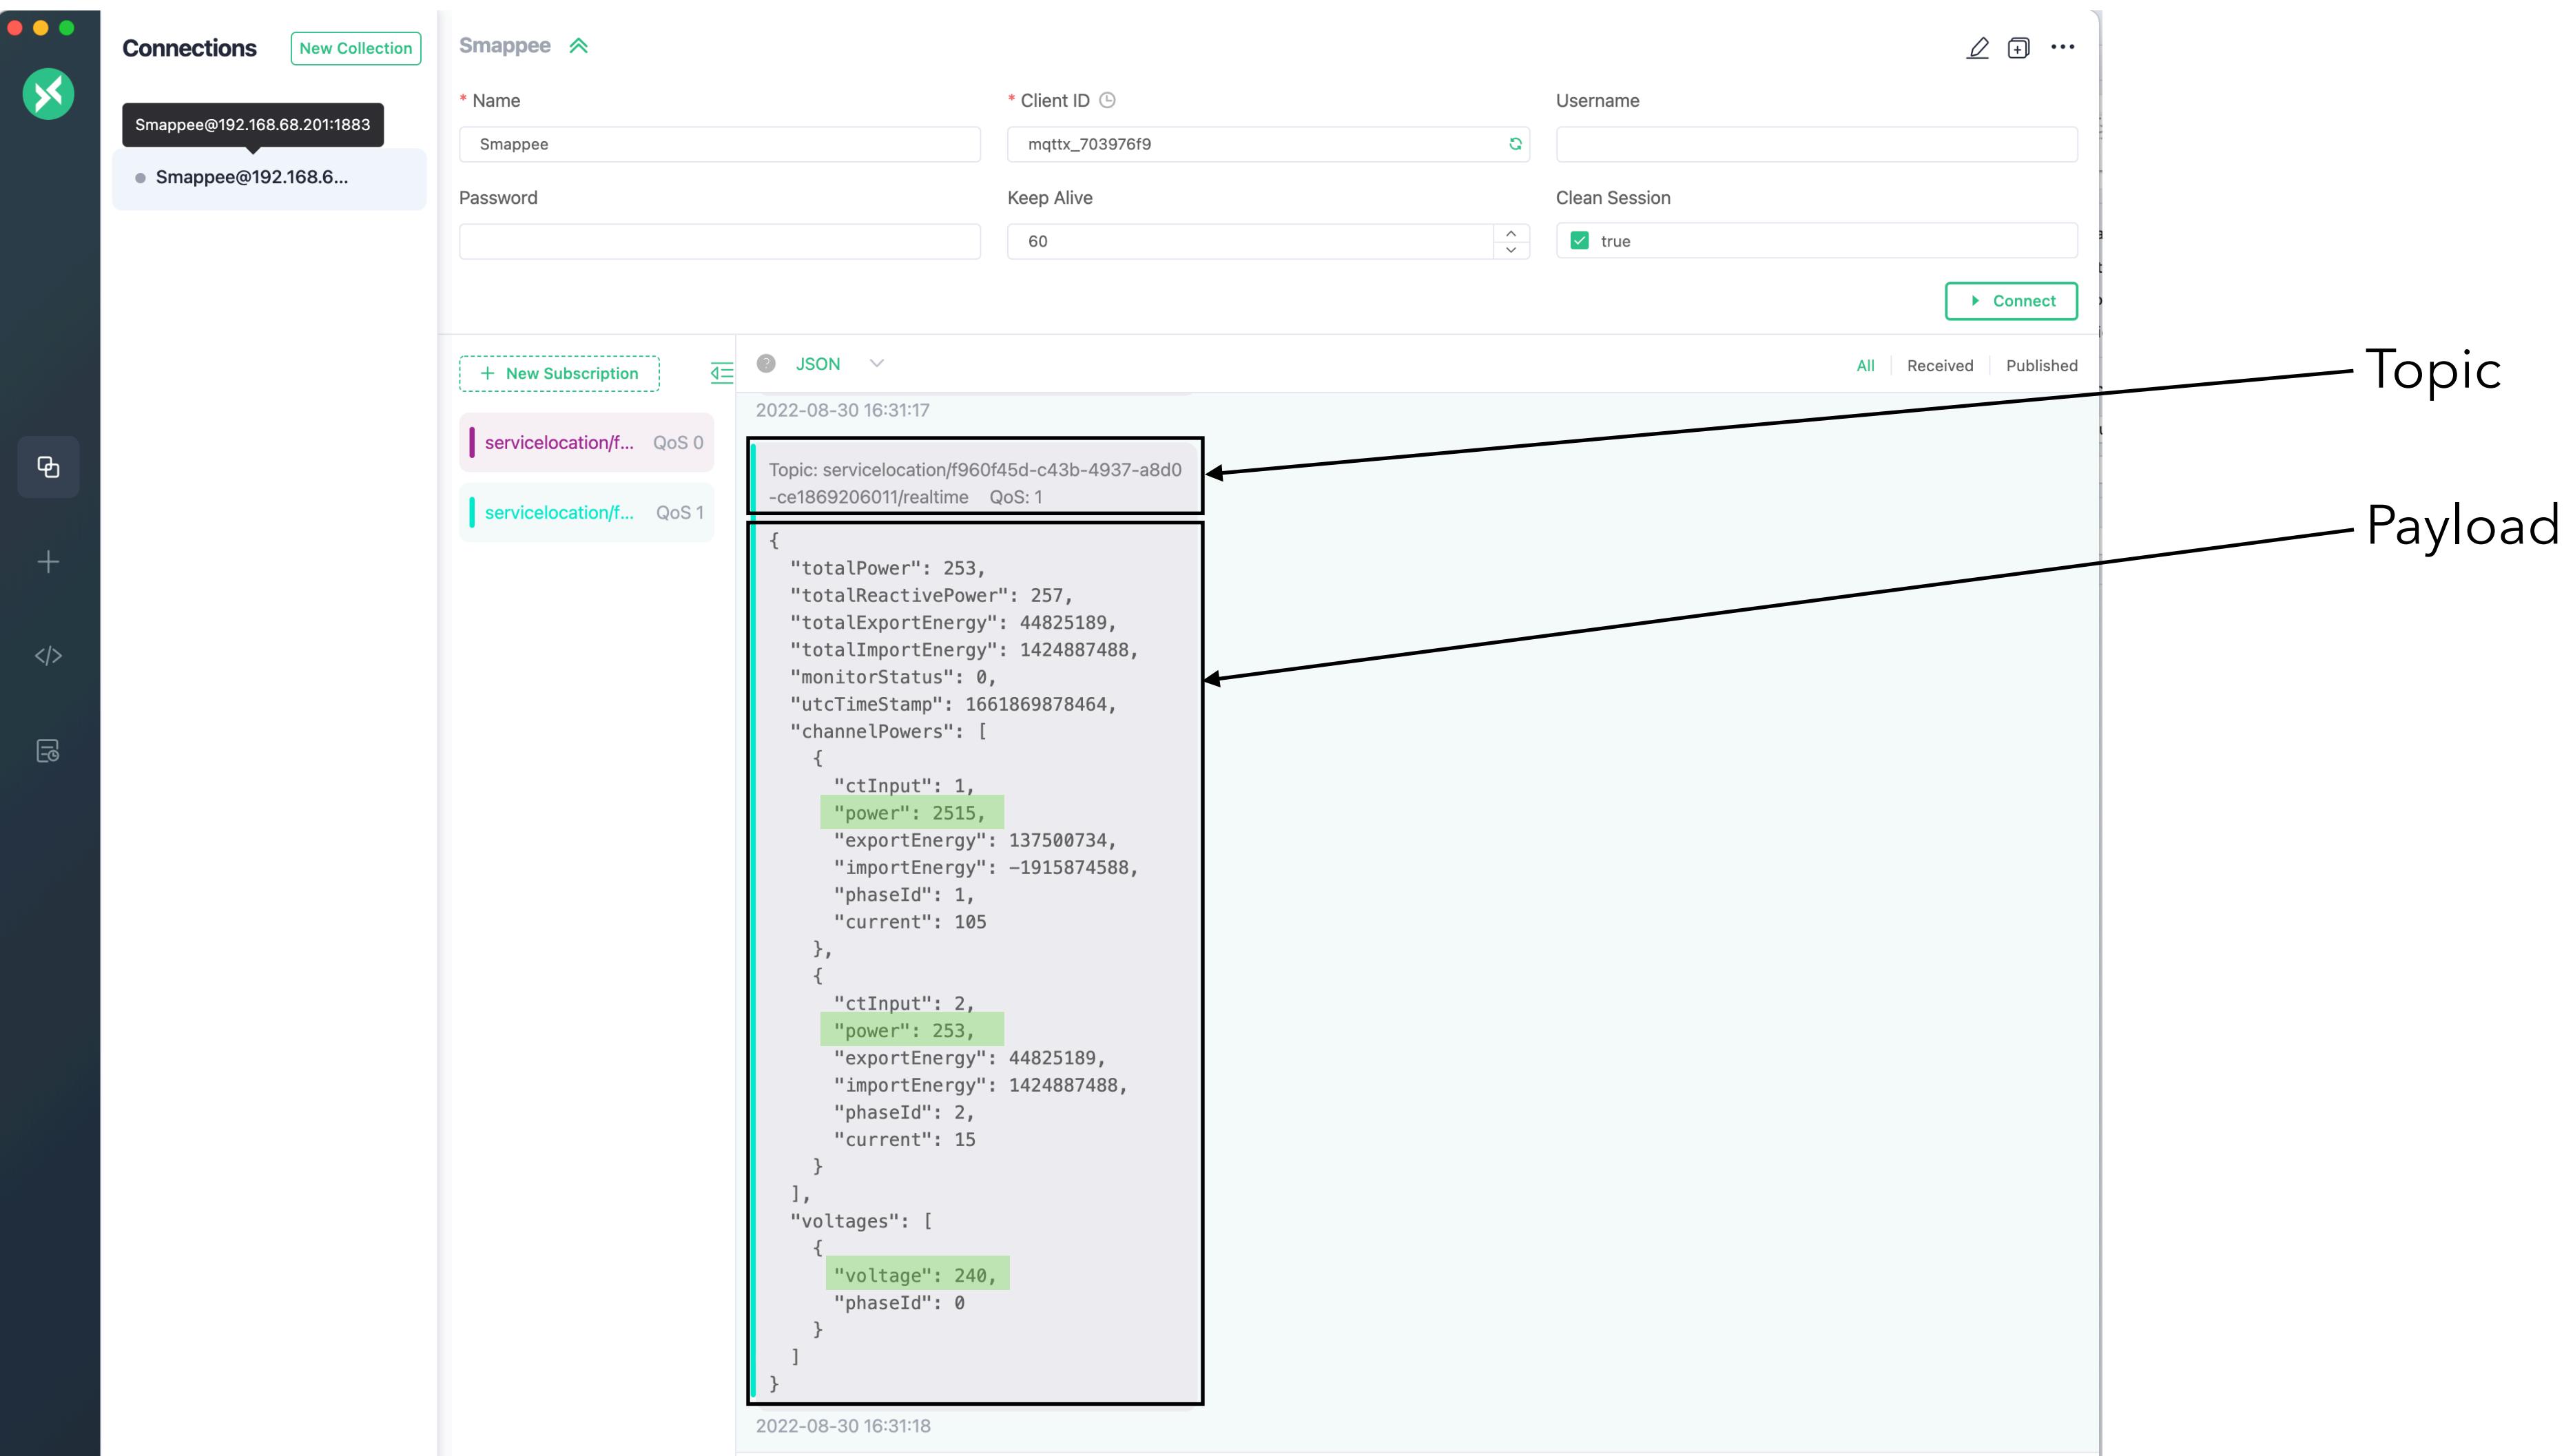
Configuring/testing MQTT

- Let's use MQTTX to find out what data our Smappee sends



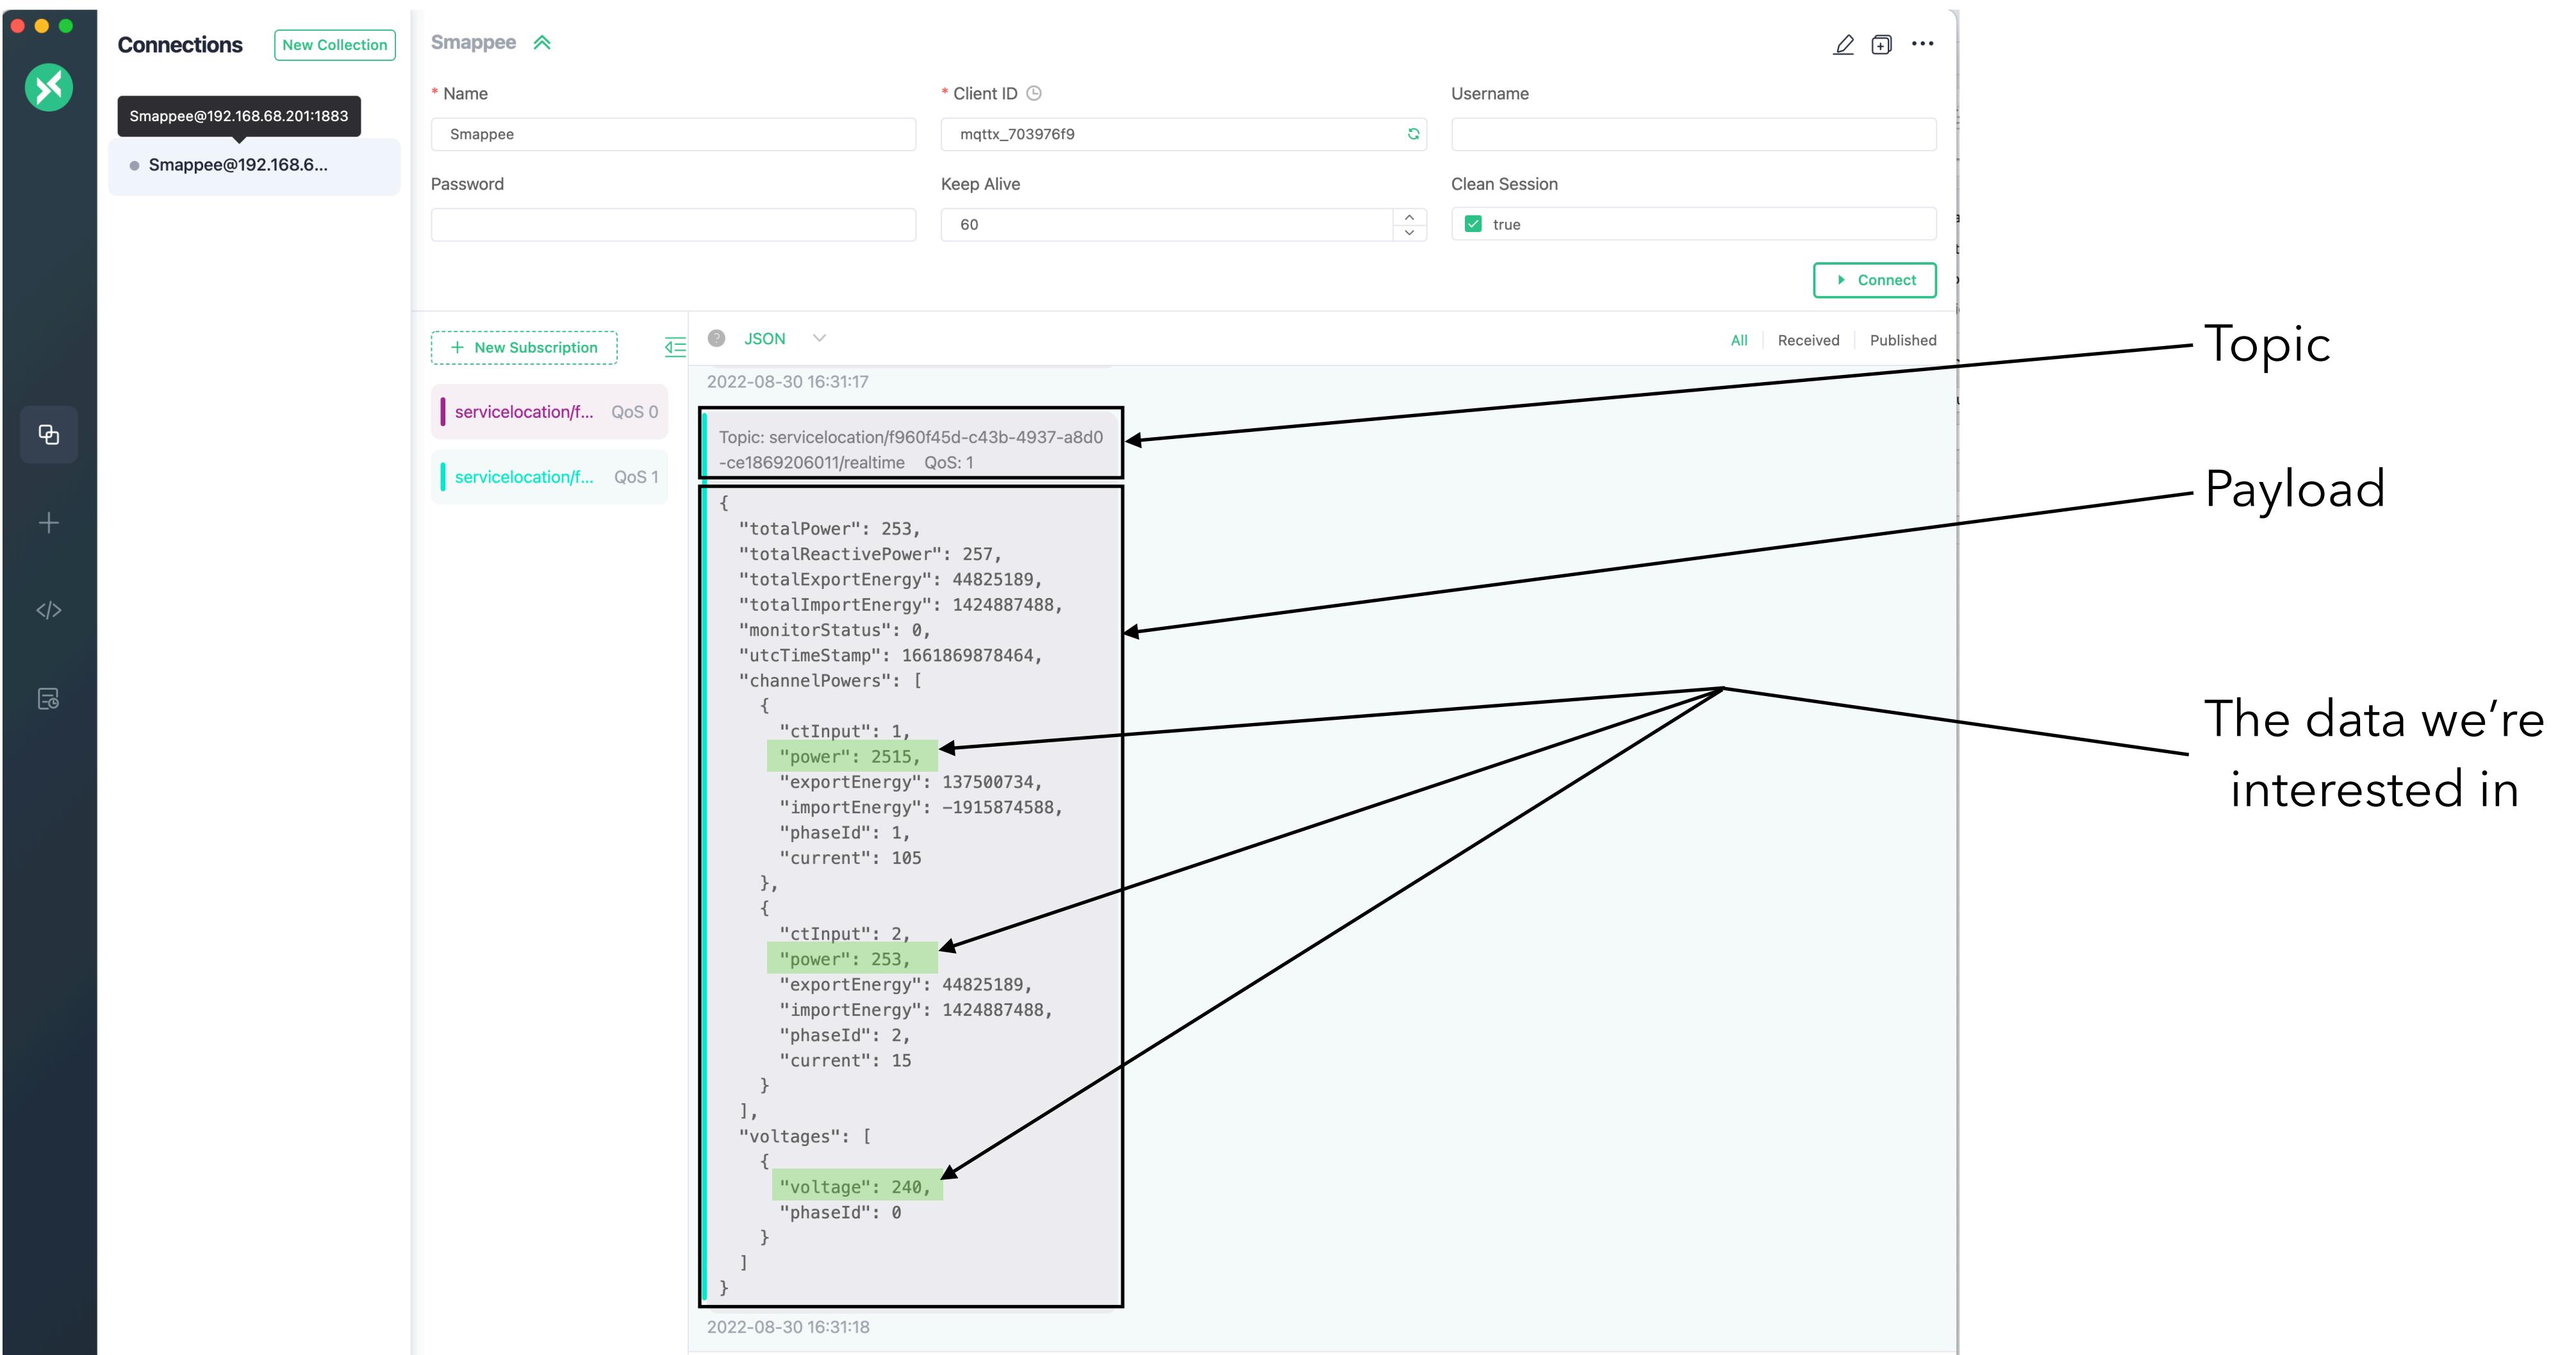
Configuring/testing MQTT

- Let's use MQTTX to find out what data our Smappee sends



Configuring/testing MQTT

- Let's use MQTTX to find out what data our Smappee sends



Telegraf configuration

- Telegraf allows to select and map data fields, aggregate them, and write them out
- This is done via configuration of [Telegraf plugins](#)

```
servicelocation/f960f45d-c43b-4937-a8d0-ce1869206011/realtime

{
  "totalPower": 253,
  "totalReactivePower": 257,
  "totalExportEnergy": 44825189,
  "totalImportEnergy": 1424887488,
  "monitorStatus": 0,
  "utcTimeStamp": 1661869878464,
  "channelPowers": [
    {
      "ctInput": 1,
      "power": 2515,
      "exportEnergy": 137500734,
      "importEnergy": -1915874588,
      "phaseId": 1,
      "current": 105
    },
    {
      "ctInput": 2,
      "power": 253,
      "exportEnergy": 44825189,
      "importEnergy": 1424887488,
      "phaseId": 2,
      "current": 15
    }
  ],
  "voltages": [
    {
      "voltage": 240,
      "phaseId": 0
    }
  ]
}
```

Telegraf configuration

- We will configure these plugins:
 - [mqtt_consumer](#) input
 - [json_v2](#) parser
 - [regex](#) processor
 - [basicstats](#) aggregator
 - [influxdb](#) output
- Let's build the configuration one step at a time

```
servicelocation/f960f45d-c43b-4937-a8d0-ce1869206011/realtim
{
  "totalPower": 253,
  "totalReactivePower": 257,
  "totalExportEnergy": 44825189,
  "totalImportEnergy": 1424887488,
  "monitorStatus": 0,
  "utcTimeStamp": 1661869878464,
  "channelPowers": [
    {
      "ctInput": 1,
      "power": 2515,
      "exportEnergy": 137500734,
      "importEnergy": -1915874588,
      "phaseId": 1,
      "current": 105
    },
    {
      "ctInput": 2,
      "power": 253,
      "exportEnergy": 44825189,
      "importEnergy": 1424887488,
      "phaseId": 2,
      "current": 15
    }
  ],
  "voltages": [
    {
      "voltage": 240,
      "phaseId": 0
    }
  ]
}
```

Telegraf configuration - mqtt_consumer input

- Telegraf plugins share a common configuration space: be precise!
- Use `name_override` and `name_pass` to wire things together
- We start with the `mqtt_consumer` input plugin:

```
[ [inputs.mqtt_consumer]
  alias = "smappee-test"
  name_override = "smappee-data-test"
  servers = ["tcp://192.168.68.201:1883"]
  topics = [
    "servicelocation/f960f45d-c43b-4937-a8d0-ce1869206011/realtim
  ]
  # The "host" tag is irrelevant in this use case. Drop it
  tagexclude = ["host"]
  data_format = "json_v2"
```

`servicelocation/f960f45d-c43b-4937-a8d0-ce1869206011/realtim`

```
{
  "totalPower": 253,
  "totalReactivePower": 257,
  "totalExportEnergy": 44825189,
  "totalImportEnergy": 1424887488,
  "monitorStatus": 0,
  "utcTimeStamp": 1661869878464,
  "channelPowers": [
    {
      "ctInput": 1,
      "power": 2515,
      "exportEnergy": 137500734,
      "importEnergy": -1915874588,
      "phaseId": 1,
      "current": 105
    },
    {
      "ctInput": 2,
      "power": 253,
      "exportEnergy": 44825189,
      "importEnergy": 1424887488,
      "phaseId": 2,
      "current": 15
    }
  ],
  "voltages": [
    {
      "voltage": 240,
      "phaseId": 0
    }
  ]
}
```

Telegraf configuration - Field selection/naming

- Next we use the json_v2 parser
- Use the [gjson.dev](#) playground to your advantage

```

[[inputs.mqtt_consumer]]
alias = "smappee-test"
name_override = "smappee-data-test"
servers = ["tcp://192.168.68.201:1883"]
topics = [
    "servicelocation/f960f45d-c43b-4937-a8d0-ce1869206011/realtime"
]
# The "host" tag is irrelevant in this use case. Drop it
tagexclude = ["host"]
data_format = "json_v2"
[[inputs.mqtt_consumer.json_v2]]
  [[inputs.mqtt_consumer.json_v2.field]]
    path = "channelPowers.#(ctInput==2).power"
    rename = "consumption"
    type = "float"
  [[inputs.mqtt_consumer.json_v2.field]]
    path = "channelPowers.#(ctInput==1).power"
    rename = "solar"
    type = "float"
  [[inputs.mqtt_consumer.json_v2.field]]
    path = "voltages.#.voltage"
    type = "float"

```

servicelocation/f960f45d-c43b-4937-a8d0-ce1869206011/realtime

```
{
  "totalPower": 253,
  "totalReactivePower": 257,
  "totalExportEnergy": 44825189,
  "totalImportEnergy": 1424887488,
  "monitorStatus": 0,
  "utcTimeStamp": 1661869878464,
  "channelPowers": [
    {
      "ctInput": 1,
      "power": 2515,
      "exportEnergy": 137500734,
      "importEnergy": -1915874588,
      "phaseId": 1,
      "current": 105
    },
    {
      "ctInput": 2,
      "power": 253,
      "exportEnergy": 44825189,
      "importEnergy": 1424887488,
      "phaseId": 2,
      "current": 15
    }
  ],
  "voltages": [
    {
      "voltage": 240,
      "phaseId": 0
    }
  ]
}
```

Telegraf configuration - Topic remapping

- The topic is a bit clunky for our purposes
- Let's rename using the `regex` processor
 - Note the utilisation of the `name_pass` directive

```
[[processors.regex]]
  namepass = ["smappee-data-test"]
  [[processors.regex.tags]]
    key = "topic"
    pattern = ".*/(.*)/.*"
    replacement = "smappee/${1}"
```

servicelocation/f960f45d-c43b-4937-a8d0-ce1869206011/realtimedata

```
{
  "totalPower": 253,
  "totalReactivePower": 257,
  "totalExportEnergy": 44825189,
  "totalImportEnergy": 1424887488,
  "monitorStatus": 0,
  "utcTimeStamp": 1661869878464,
  "channelPowers": [
    {
      "ctInput": 1,
      "power": 2515,
      "exportEnergy": 137500734,
      "importEnergy": -1915874588,
      "phaseId": 1,
      "current": 105
    },
    {
      "ctInput": 2,
      "power": 253,
      "exportEnergy": 44825189,
      "importEnergy": 1424887488,
      "phaseId": 2,
      "current": 15
    }
  ],
  "voltages": [
    {
      "voltage": 240,
      "phaseId": 0
    }
  ]
}
```

Telegraf configuration - Topic remapping

- The topic is a bit clunky for our purposes
- Let's rename using the regex processor
 - Note the utilisation of the name_pass directive

```
[ [processors.regex]
  namepass = ["smappee-data-test"]
  [processors.regex.tags]
    key = "topic"
    pattern = ".*/(.*)/.%"
    replacement = "smappee/${1}"
```

servicelocation/f960f45d-c43b-4937-a8d0-ce1869206011 realtime



smappee/f960f45d-c43b-4937-a8d0-ce1869206011

servicelocation/f960f45d-c43b-4937-a8d0-ce1869206011 realtime

```
{
  "totalPower": 253,
  "totalReactivePower": 257,
  "totalExportEnergy": 44825189,
  "totalImportEnergy": 1424887488,
  "monitorStatus": 0,
  "utcTimeStamp": 1661869878464,
  "channelPowers": [
    {
      "ctInput": 1,
      "power": 2515,
      "exportEnergy": 137500734,
      "importEnergy": -1915874588,
      "phaseId": 1,
      "current": 105
    },
    {
      "ctInput": 2,
      "power": 253,
      "exportEnergy": 44825189,
      "importEnergy": 1424887488,
      "phaseId": 2,
      "current": 15
    }
  ],
  "voltages": [
    {
      "voltage": 240,
      "phaseId": 0
    }
  ]
}
```

Telegraf configuration - Aggregation

- Data arriving at 1 sample/s
- Aggregate to 1 minute average
 - Drop the original data

```

[[aggregators.basicstats]]
  namepass = ["smappee-data-test"]
  ## The period on which to flush & clear the aggregator.
  period = "60s"

  ## If true, the original metric will be dropped by the
  ## aggregator and will not get sent to the output plugins.
  drop_original = true

  ## Configures which basic stats to push as fields
  stats = ["mean"]

```

servicelocation/f960f45d-c43b-4937-a8d0-ce1869206011/realtim

```
{
  "totalPower": 253,
  "totalReactivePower": 257,
  "totalExportEnergy": 44825189,
  "totalImportEnergy": 1424887488,
  "monitorStatus": 0,
  "utcTimeStamp": 1661869878464,
  "channelPowers": [
    {
      "ctInput": 1,
      "power": 2515,
      "exportEnergy": 137500734,
      "importEnergy": -1915874588,
      "phaseId": 1,
      "current": 105
    },
    {
      "ctInput": 2,
      "power": 253,
      "exportEnergy": 44825189,
      "importEnergy": 1424887488,
      "phaseId": 2,
      "current": 15
    }
  ],
  "voltages": [
    {
      "voltage": 240,
      "phaseId": 0
    }
  ]
}
```

Telegraf configuration - Output to InfluxDB

- Persisting data to InfluxDB is a breeze

```
[[outputs.influxdb]]
namepass = ["smappee-data-test"]
alias = "smappee-out"
urls = ["http://192.168.68.201:8086"]
database = "smappee_test"
username = "demo"
password = "demo"
```

servicelocation/f960f45d-c43b-4937-a8d0-ce1869206011/realtimedata

```
{
  "totalPower": 253,
  "totalReactivePower": 257,
  "totalExportEnergy": 44825189,
  "totalImportEnergy": 1424887488,
  "monitorStatus": 0,
  "utcTimeStamp": 1661869878464,
  "channelPowers": [
    {
      "ctInput": 1,
      "power": 2515,
      "exportEnergy": 137500734,
      "importEnergy": -1915874588,
      "phaseId": 1,
      "current": 105
    },
    {
      "ctInput": 2,
      "power": 253,
      "exportEnergy": 44825189,
      "importEnergy": 1424887488,
      "phaseId": 2,
      "current": 15
    }
  ],
  "voltages": [
    {
      "voltage": 240,
      "phaseId": 0
    }
  ]
}
```

Telegraf configuration - Summary

- The complete configuration at a glance

```

[[inputs.mqtt_consumer]]
alias = "smappee-test"
name_override = "smappee-data-test"
servers = ["tcp://192.168.68.201:1883"]
topics = [
    "servicelocation/f960f45d-c43b-4937-a8d0-ce1869206011/realtim"
]

tagexclude = ["host"]
data_format = "json_v2"

[[inputs.mqtt_consumer.json_v2]]

  [[inputs.mqtt_consumer.json_v2.field]]
    path = "channelPowers.#(ctInput==2).power"
    rename = "consumption"
    type = "float"

  [[inputs.mqtt_consumer.json_v2.field]]
    path = "channelPowers.#(ctInput==1).power"
    rename = "solar"
    type = "float"

  [[inputs.mqtt_consumer.json_v2.field]]
    path = "voltages.#.voltage"
    type = "float"

[[processors.regex]]
namepass = ["smappee-data-test"]
[[processors.regex.tags]]
key = "topic"
pattern = ".*/(.*)/.*"
replacement = "smappee/${1}"

[[aggregators.basicstats]]
namepass = ["smappee-data-test"]
## The period on which to flush & clear the aggregator.
period = "60s"
drop_original = true
stats = ["mean"]

[[outputs.influxdb]]
namepass = ["smappee-data-test"]
alias = "smappee-out"
urls = ["http://192.168.68.201:8086"]
database = "smappee_test"
username = "demo"
password = "demo"

```

InfluxDB configuration

- Using InfluxDB version 1
- The configuration is extremely simple using the InfluxDB CLI:
 - Create the database(s)
 - Create a user with write privileges

```
akkapi@node-1:~$ influx
Connected to http://localhost:8086 version 1.8.10
InfluxDB shell version: 1.8.10
> create database smappee_test
> create user demo with password 'demo' with all privileges;
```

- ...

InfluxDB version 1 versus version 2

- InfluxDB version 2 ...
 - **TICK** stack in a single binary:
 - Telegraf
 - InfluxDB
 - Chronograf (Dashboards, Admin, Alerting)
 - Kapacitor (Real-time streaming data processing engine)
 - Has much better security than version 1, authentication baked in everywhere
 - Has Flux Query (and still supports InfluxQL)
 - InfluxDB tasks for downsampling data
 - Has a great upgrade story as I experienced myself
 - ...

Ubuntu - takeaways

- Using cloud-init, a system can be provisioned from scratch in about 20 minutes
- Ubuntu 20.04 is not the most recent version
 - Cannot directly boot off an external SSD. An SD still has to be present

```
akkapi@node-1:~$ mount -t ext4,vfat  
/dev/sda2 on / type ext4 (rw,relatime)  
/dev/sdal on /boot/firmware type vfat (rw,relatime,fmask=0022,dmask=0022,codepage=437,iocharset=ascii,shortname=mixed,errors=remount-ro)  
akkapi@node-1:~$ _
```

- Flashing an SD takes ±5 minutes
- Flashing an SSD takes ±18 seconds

cloud-init - takeaways

- [cloud-init](#) is a method for cross-platform instance initialisation
- Can be used for bare-metal installations
- Performs user creation, execute custom scripts, install packages, create files, ...
- Has rather poor docs, but this has improved a lot recently
- Read the [Module Reference](#)
- Supports [jinja](#) template rendering:
 - cloud-init meta-data file contains key/value pairs
 - keys used in user-data will be dereferenced by jinja

```
instance_id: cloud-image
influx_db1: smappee_test
influx_user: demo
influx_passwd: demo
```

meta-data

```
runcmd:
- |
  sudo -uakkapi influx <<ENDCFGINFLUX
  create database {{ ds.meta_data.influx_db1 }};
  create user {{ ds.meta_data.influx_user }} with password '{{ ds.meta_data.influx_passwd }}' with all privileges;
ENDCFGINFLUX
```

user-data

cloud-init - jinja

- Your meta-data key-value pairs may not be where you expect them to be

```
akkapi@node-4:~$ cloud-init query --all
{
    "_beta_keys": [
        "subplatform"
    ],
    "availability_zone": null,
    "base64_encoded_keys": [],
    "cloud_name": "unknown",
    "distro": "ubuntu",
    "distro_release": "focal",
    "distro_version": "20.04",
    "ds": {
        "_doc": "EXPERIMENTAL: The structure and format of content scoped under the 'ds' key may change in subsequent releases of cloud-init.",
        "meta_data": {
            "dsmode": "net",
            "influx_db1": "smappee_test",
            "influx_passwd": "demo",
            "influx_user": "demo",
            "instance_id": "cloud-image"
        }
    },
    "instance_id": "nocloud",
    <elided>
```

cloud-init - jinja

- Flashing

```
$ flash -n node-1 -j -m cloud-init/meta-data -u cloud-init/akka-pi-os-dhcp-64-ssd.yml \
https://cdimage.ubuntu.com/releases/20.04/release/ubuntu-20.04.5-preinstalled-server-arm64+raspi.img.xz
```

What's next?

- Explore InfluxDB version 2
- Add a new Smappee device with 12 channels to my home setup
- More Grafana dashboards
- More aggregations
- Alerts
- Automatic equipment control based on data (for example, an electrical water boiler)

Conclusions

- Learned a ton about InfluxDB, Telegraf, EMQX, Grafana
- Documentation is a weak point of many of the OSS projects, but things are improving
- Implemented the monitoring system that:
 - ✓ avoided programming
 - ✓ uses Open Source Software only
 - ✓ captures all knowledge through installation and configuration automation
 - ✓ automates the installation of the OS and required software





Questions?